



Introduction to OpenMP

R. Bader (LRZ)

G. Hager (RRZE)

V. Weinberg (LRZ)

1. Increase performance / throughput of CPU core

- a) Reduce cycle time, i.e. increase clock speed (Moore)
- b) Increase throughput, i.e. superscalar + SIMD

2. Improve data access time

- a) Increase cache size
- b) Improve main memory access (bandwidth & latency)




3. Use parallel computing (shared memory)

- a) Requires shared-memory parallel programming
- b) Shared/separate caches
- c) Possible memory access bottlenecks

4. Use parallel computing (distributed memory)


“Cluster” of computers tightly connected

- a) Almost unlimited scaling of memory and performance
- b) Distributed-memory parallel programming



5. Use an accelerator with your compute node

- a) Requires offload of program regions (semantics may be limited)
- b) Host and accelerator memory are connected, but separate



(Improvements are under way)

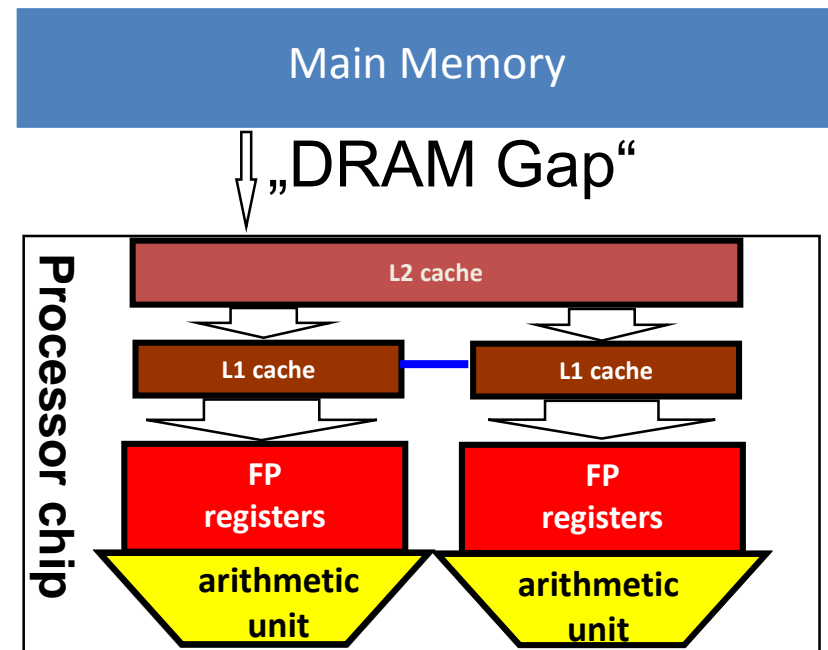

- c) Programming complexity is higher than for shared memory systems („heterogeneous parallel computing“)

It is not a faster CPU – it is a **parallel computer on a chip**.


Put multiple processors (“cores”) on a chip which share resources (example shows a dual core that shares L2 cache and memory bandwidth)

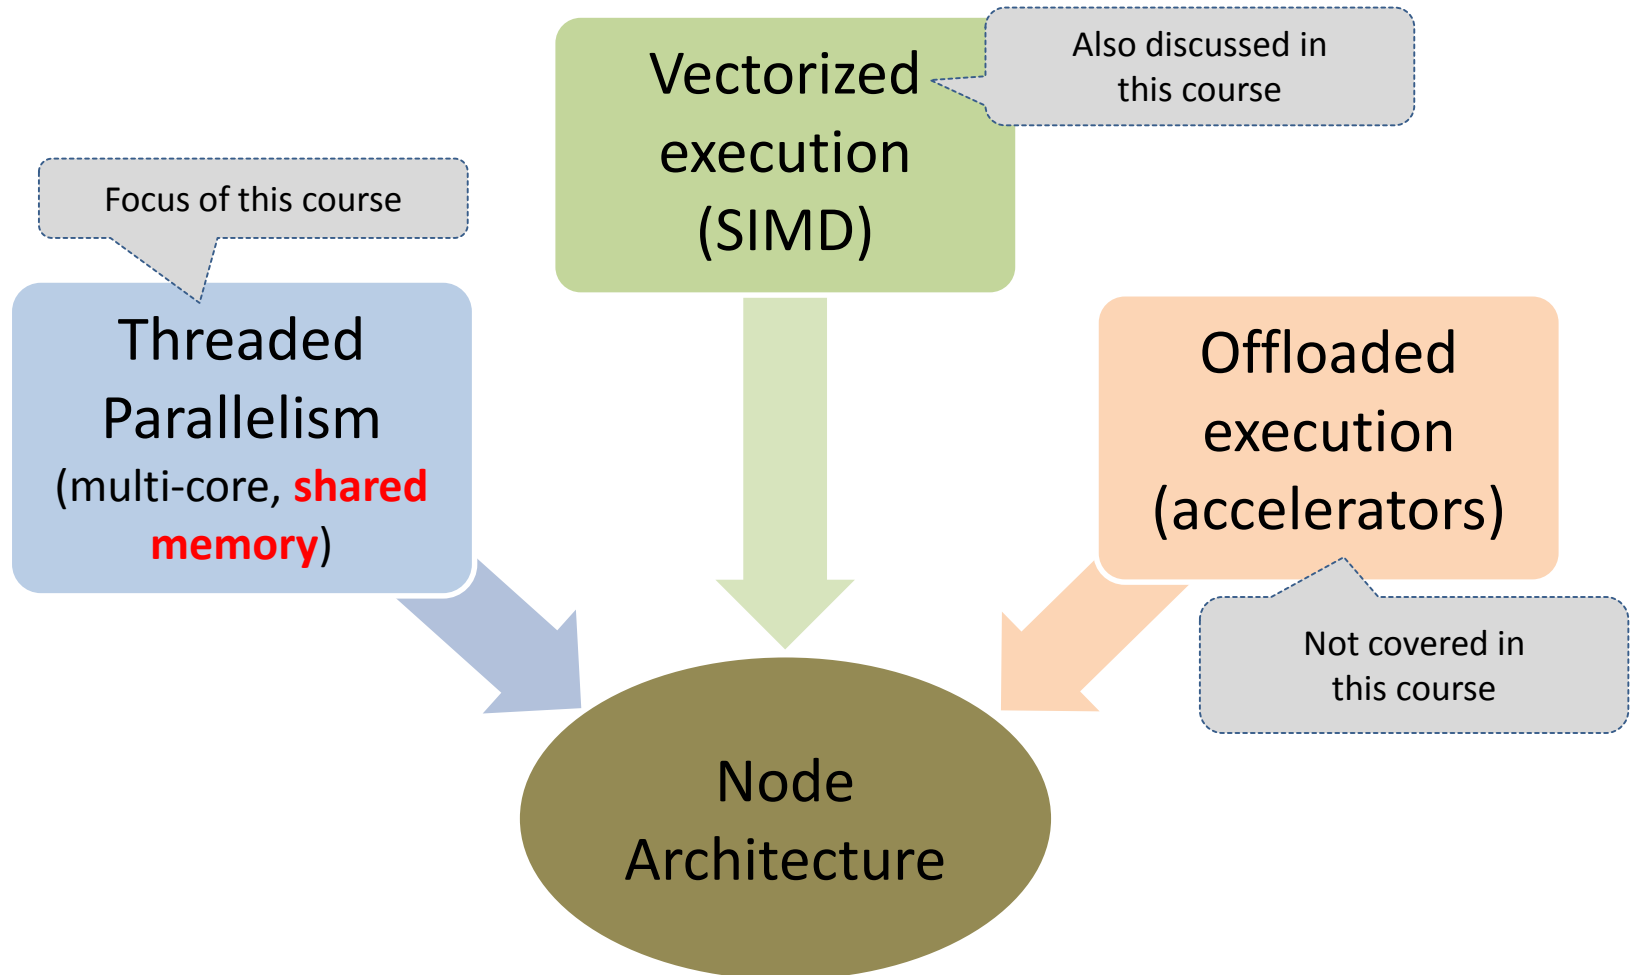
Efficient use of all cores for a single application → programmer

Intel Xeon (Woodcrest)



- Option 1 a) is not feasible any more, option 2 only in small increments





■ Syntactic portability

- Directives / pragmas
- Conditional compilation permits to mask API calls

■ Semantic portability

- Standardized across platforms
→ safe-to-use interface
- Unsupported/unavailable hardware features → irrelevant directives will be ignored
(you might need a special compiler for your devices ...)

■ Performance portability

- Unfortunately performance is not necessarily portable
- Has traditionally been a problem
(partly due to differences in hardware/architectural properties)

Are semantics for sequential execution retained?

- yes, due to directive concept
- programmer may **choose** not to

Do memory accesses occur in the same order?

- no, due to **relaxed** memory consistency (performance feature!)

Are the same numeric results obtained for parallel execution?

- **no associativity** for model number operations
- parallel execution might reorder operations
(programmer may need to enforce ordering for reproducibility and/or numeric stability)

- **Responsible body:** OpenMP Architecture Review Board

- Published OpenMP **4.5** in November 2015
- Development continues

History of OpenMP
starts in 1997

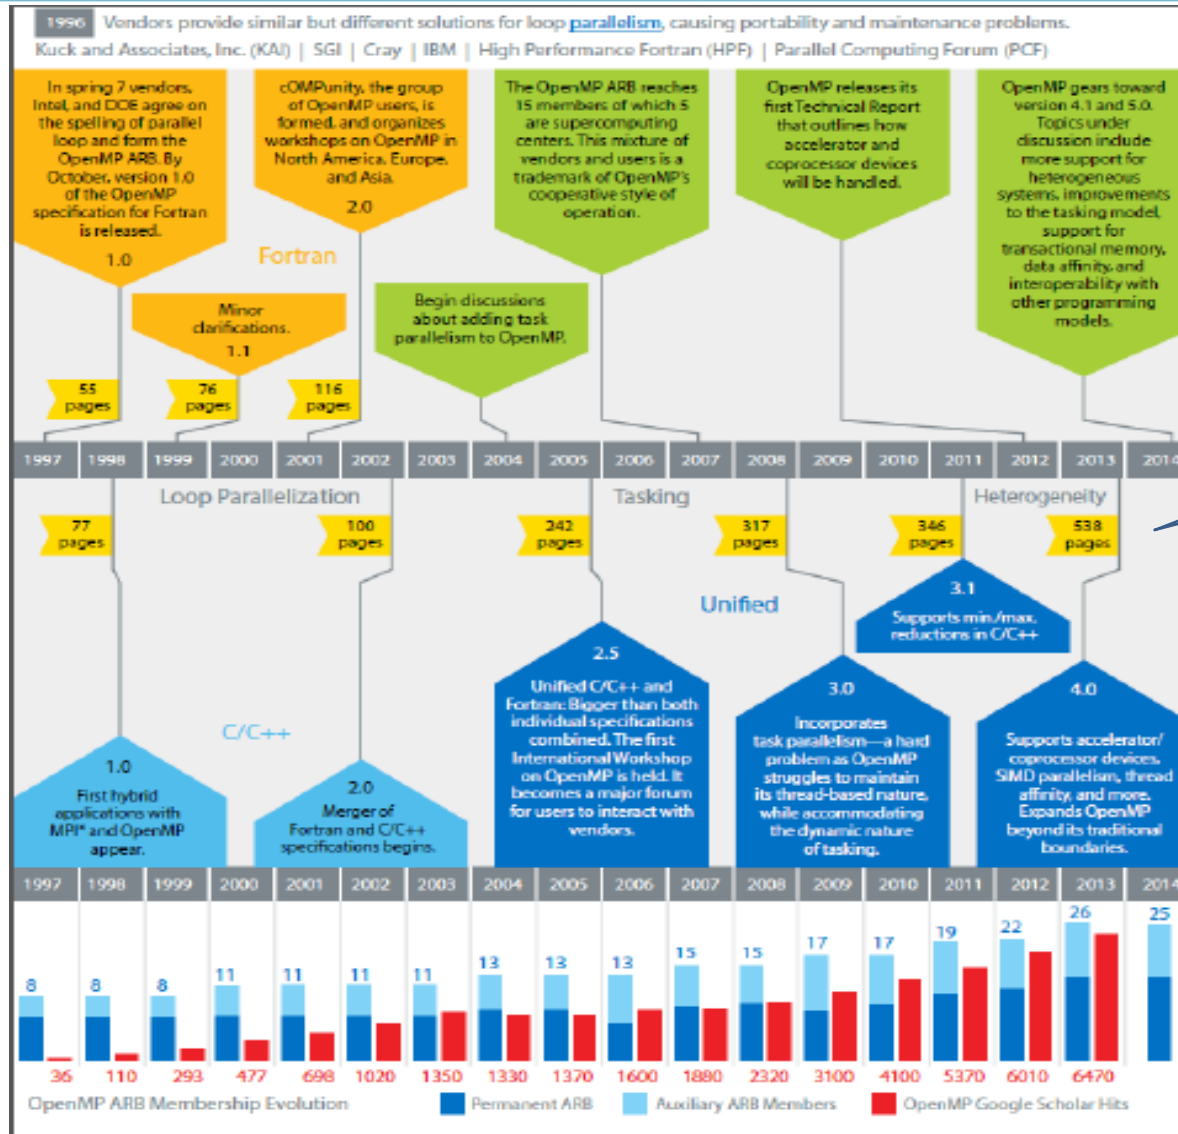
- **Base languages**

- Fortran (77, 95, 2003)
- C, C++
- (Java is not a base language)

Fortran and C examples
will be displayed

- **Resources:**

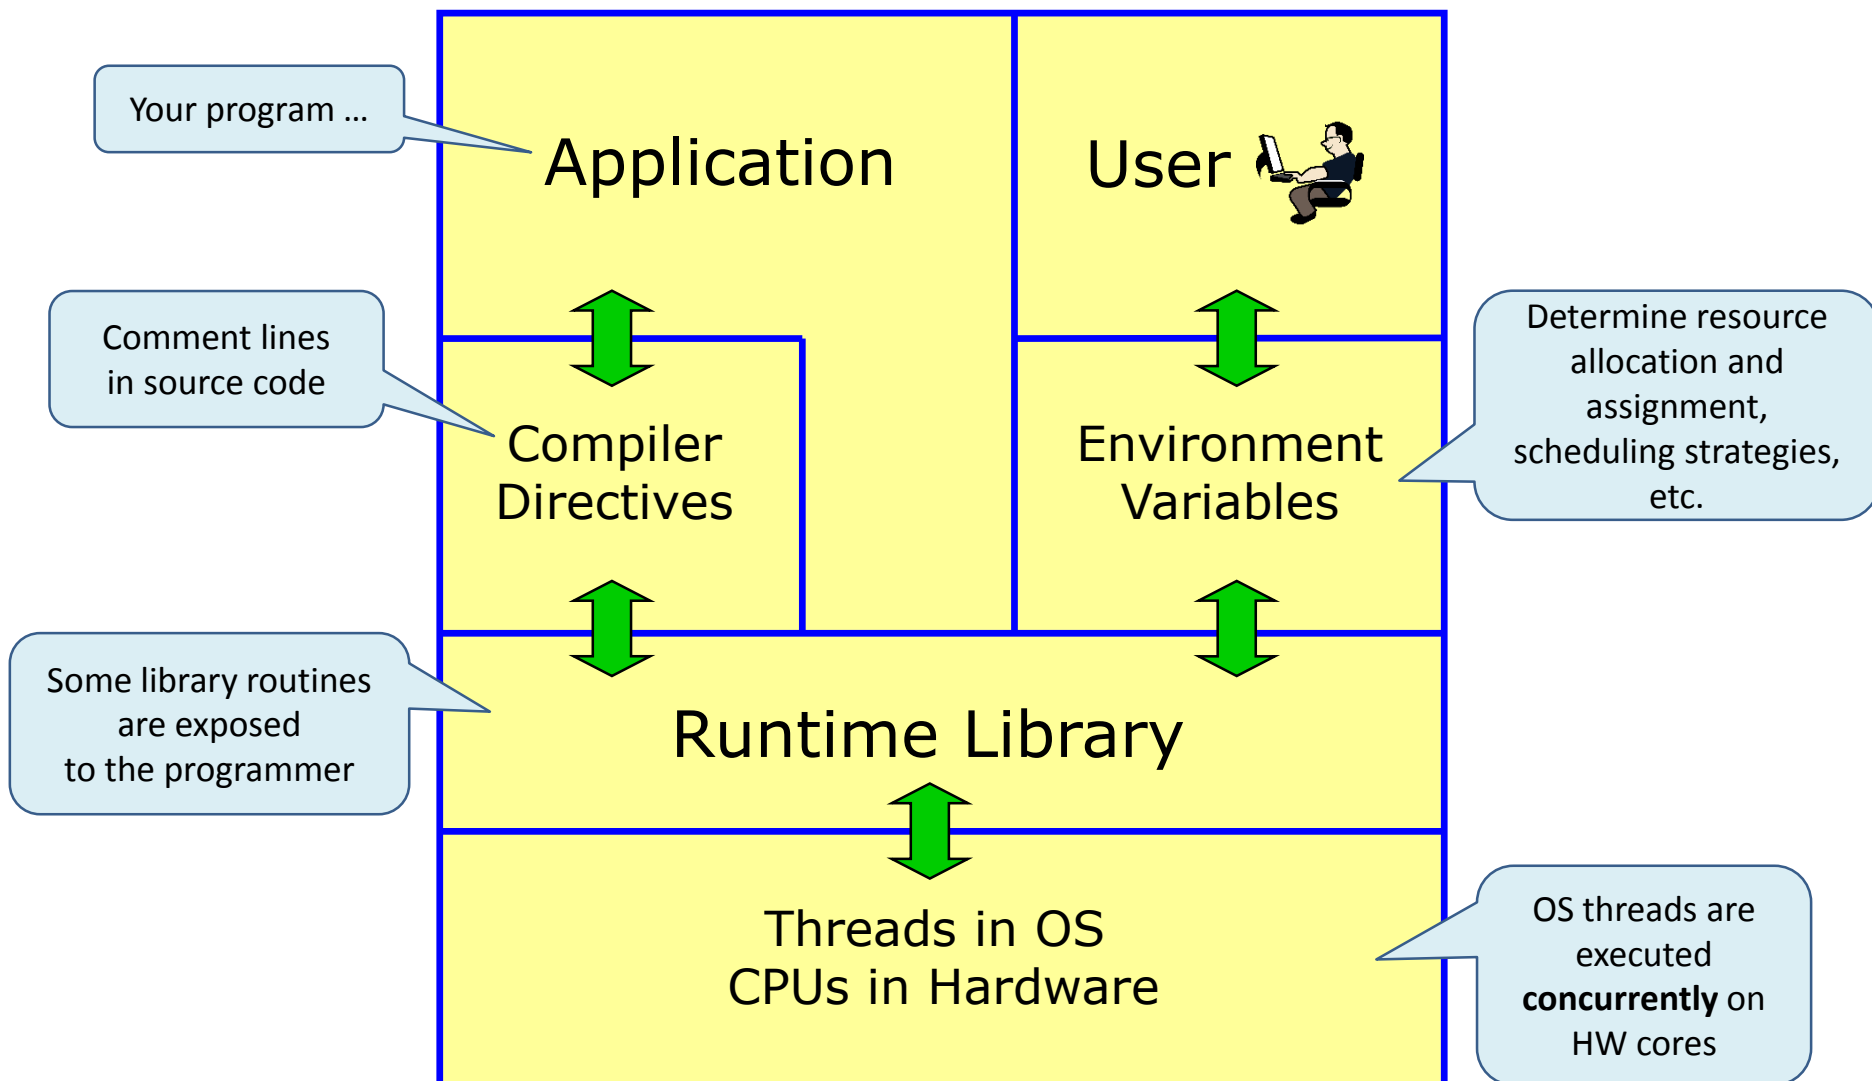
- <http://www.openmp.org> (including standard documents)
- <http://www.compunity.org>



Note the increase in the standard's size

Course Target:

Learn the most useful and therefore most commonly used features of OpenMP



Fortran

```
program
  use m
  implicit none

  call f()

end program

module m
  implicit none
contains
  subroutine f()
    print *, 'Hello'
  end subroutine
end module
```

Aim is to execute
f() in parallel

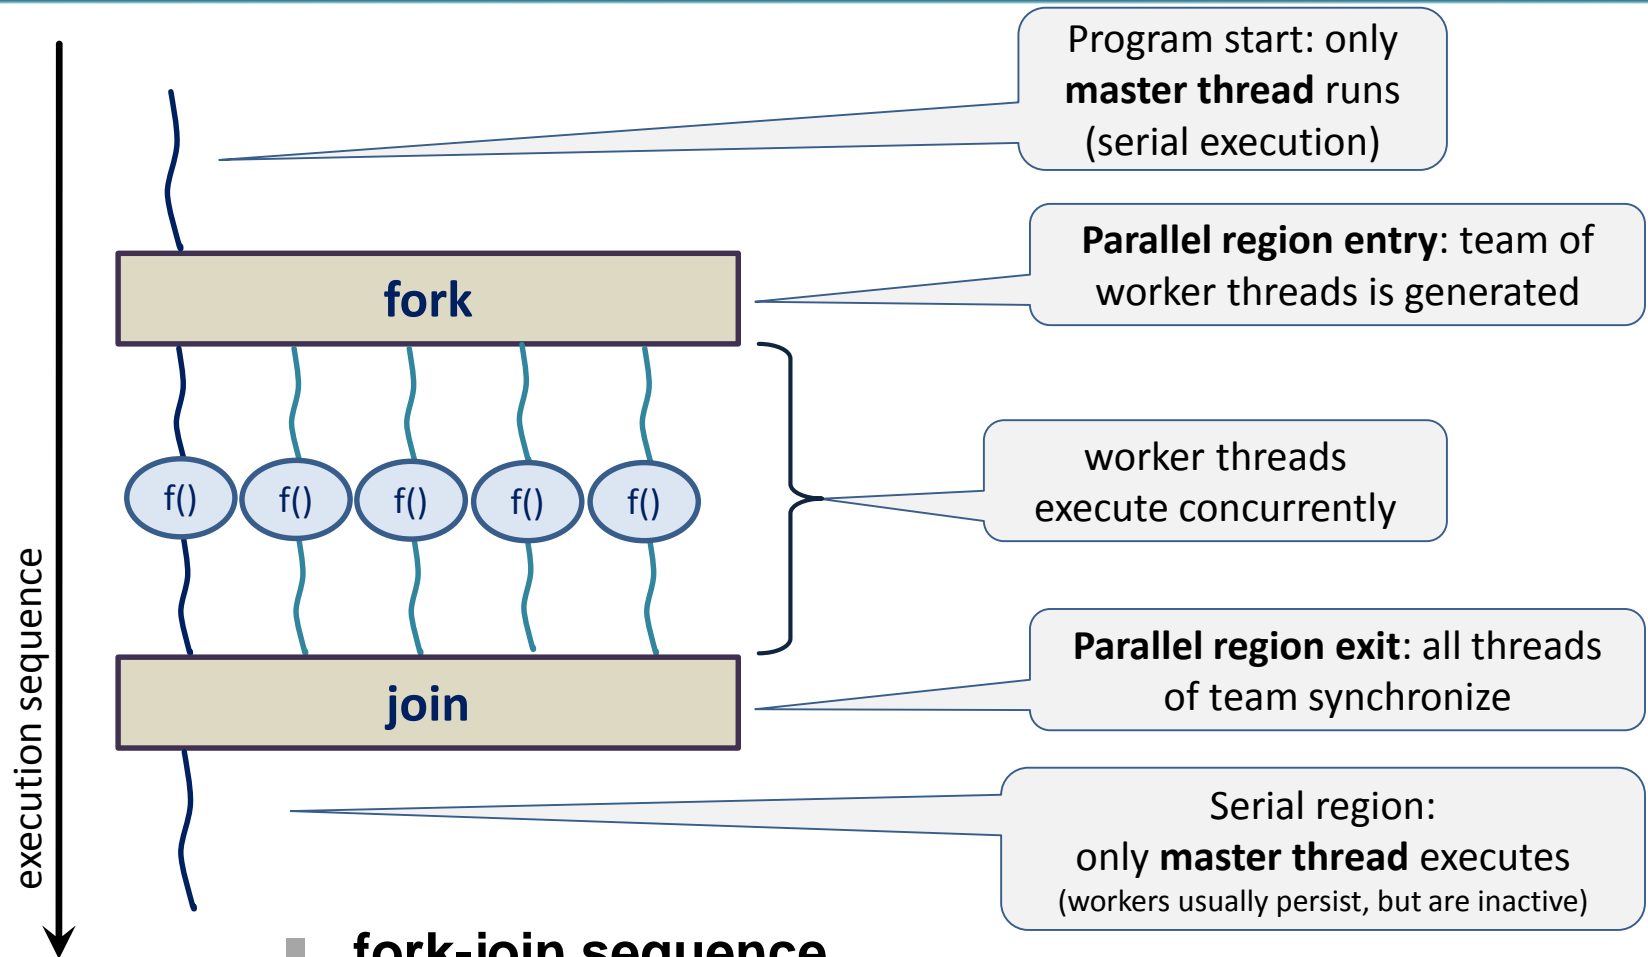
C

```
#include <stdio.h>
int main() {

  f();

  return 0;
}

void f() {
  printf("Hello\n");
}
```



- **fork-join sequence**
 - can repeat, with differing thread counts

Fortran

```

program
  use m
  implicit none
!$omp parallel
  call f()
!$omp end parallel
end program

```

enclosed
lexical block

C

```

#include <stdio.h>
int main() {
#pragma omp parallel
{
  f();
}
return 0;
}

```

■ **General form of directives:**

```
!$omp <directive> [<clause>]
```

sentinel

```
#pragma omp <directive> [<clause>]
```

sentinel

- clauses, if present, modify a directive's semantics
- multiple clauses per directive are possible
- continuation lines are supported for long directives:

Fortran

&

C

\

Fortran

- statements between a beginning and ending directive pair

single point of entry

- GOTO into block is prohibited

single point of exit

- GOTO, RETURN, EXIT outside block are prohibited

permitted: program termination

- STOP, ERROR STOP

C / C++

- delineated by braces following a directive

- setjmp() into block is prohibited

- longjmp() and throw() outside block are prohibited

- exit()

Fortran

```
subroutine f()
!$ use omp_lib
integer :: me
me = 0
!$ me = omp_get_thread_num()
print *, 'Hello from thread ', me
end subroutine
```

OpenMP module:
explicit interfaces for API

returns an integer
(avoid implicit typing!)

!\$ indicates statement should
be compiled **conditionally**


C

```
#include <stdio.h>
#include <omp.h>
void f() {
    int me = 0;
#ifdef _OPENMP
    me = omp_get_thread_num();
#endif
    printf("Hello from thread %i\n",me);
}
```

OpenMP include file:
prototypes for API

OpenMP-specific macro for
conditional compilation

- As many independent function calls as there are threads
- Thread-individual memory management within function call
 - local variables ("me") are created in the thread-specific stack
 - malloc() or ALLOCATE create memory in the heap separately for each thread
- Private variables
 - associated with a particular thread are **inaccessible** by any other thread
 - **pro:** **safe** to use
 - **con:** **communication** is not possible (it is needed by many parallel algorithms), unnecessary replication of objects may happen.
- Thread-individual stack limit
 - control via environment variable (example: 100 MByte)



```
export OMP_STACKSIZE=100M
```

■ Classes of routines:

- Execution environment (36), **Locking (12)**, Timing (2), **Device Memory (7)**

most commonly used subset

Name	Result type	Purpose
omp_set_num_threads (int num_threads)	none	number of threads to be created for subsequent parallel region
omp_get_num_threads()	int	number of threads in currently executing region
omp_get_max_threads()	int	maximum number of threads that can be created for a subsequent parallel region
omp_get_thread_num()	int	thread number of calling thread (zero based) in currently executing region
omp_get_num_procs()	int	number of processors available
omp_get_wtime()	double	return wall clock time in seconds since some (fixed) time in the past
omp_get_wtick()	double	resolution of timer in seconds

■ **Compilation:**

Fortran

```
f90 -fopenmp -o hello.exe hello.f90
```

generic instructions ...

C

```
cc -fopenmp -o hello.exe hello.c
```

■ **Switch for OpenMP**

- specific spelling is compiler-dependent
- toggles both directives and conditional compilation
- generates threaded code and links against OpenMP run time

serial compilation may
require stub library

■ **Execution:**

```
export OMP_NUM_THREADS=4  
./hello.exe
```

by default, parallel regions
generate a team with 4 threads

■ **Output for example
program:**

```
Hello from 1  
Hello from 3  
Hello from 0  
Hello from 2
```

ordering will vary between runs
(asynchronous execution)


Now: First exercise session



Simple work sharing, Scoping of Data, and Synchronization


- **We know how to set up threading, but**
 - how can a large work item be divided up among threads?
(using the API for this works in principle, but is tedious)
 - what happens with objects that already exist before the parallel region starts?
- **Example:**
 - matrix-vector multiplication $r = M \cdot x$ i.e.

$$r_i = \sum_{j=1}^n M_{ij} x_j$$



A bunch of scalar products

- The idea is to split the work among threads



- **Note that**
 - all elements of x must be available to **all** threads
 - Matrix-Vector is often deployed iteratively $\rightarrow r$ becomes x in the next iteration \rightarrow copying of data must be possible
- **Consequence:**
 - need for variables that are accessible to **all** threads
 - \rightarrow "data sharing" is often a prerequisite for "work sharing"
 - \rightarrow a natural concept for a shared memory programming model

```
real :: s, a(200)
```

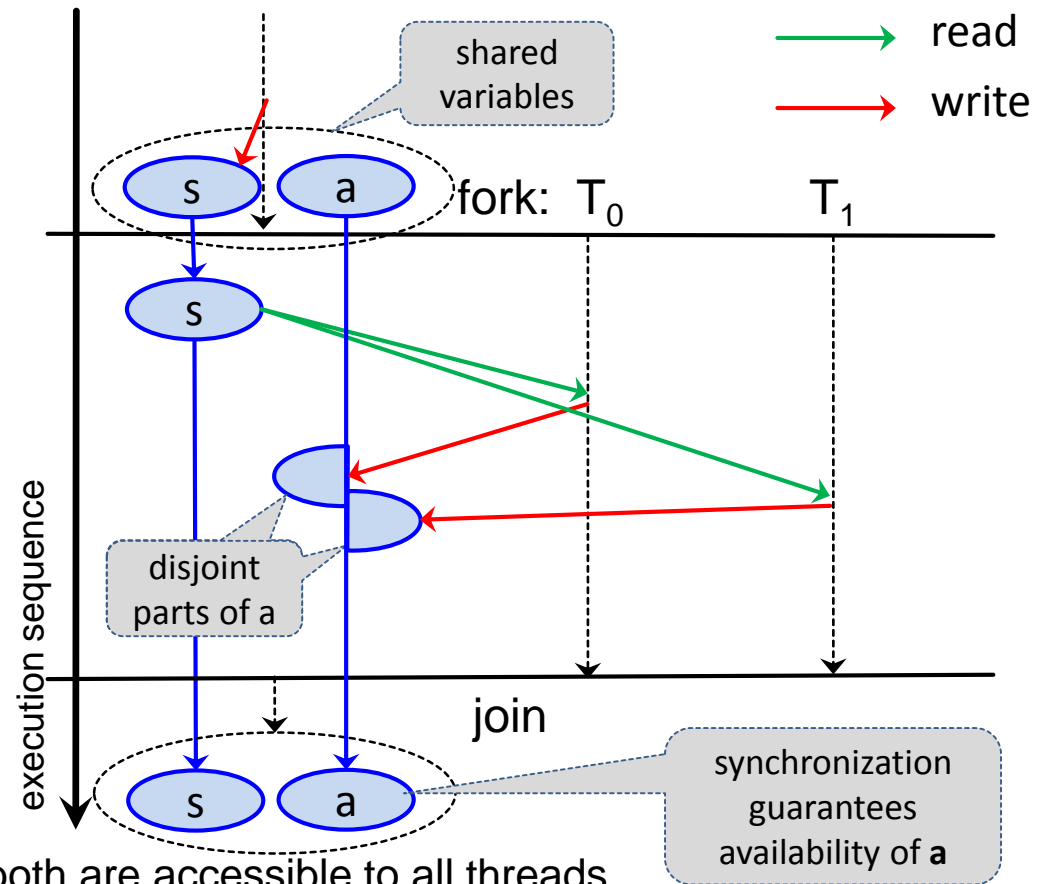
Fortran

```

s = ...
!$omp parallel shared(s,a)
  select case (me)
  case (0)
    a(1:100) = ... * s
  case (1)
    a(101:200) = ... * (-s)
  end select
!$omp end parallel

```

thread ID



■ The „shared“ clause

- implies that scalar **s** and array **a** both are accessible to all threads




Rules for concurrent accesses to a single object


- reads/writes or writes/writes by different threads are **not permitted** („data races“)

Note: updates to array **a** are OK because **disjoint parts** of object are updated


- Flow dependency ("read after write", RAW):



- Anti-dependency ("write after read", WAR):



- Output dependency ("write after write", WAW):



C

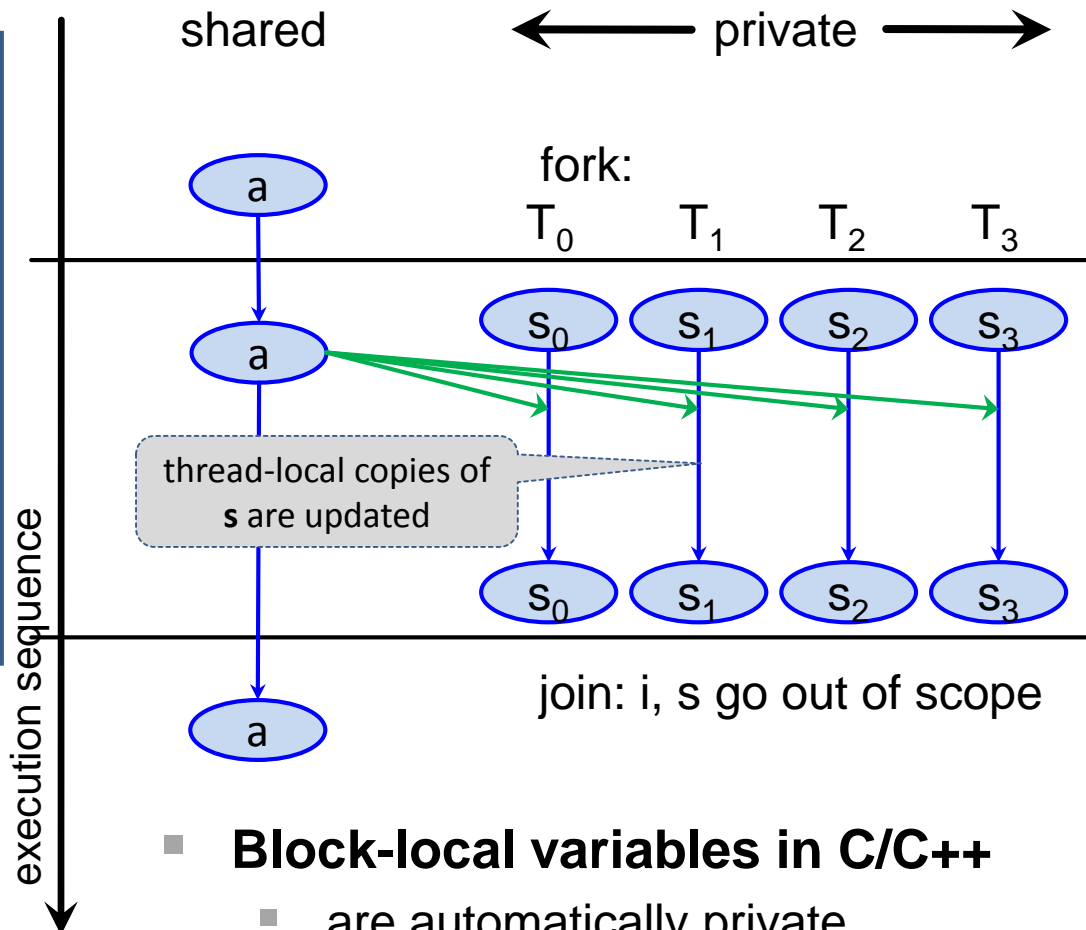
```

a[k] = ...;
#pragma omp parallel \
    shared(a)
{ int i; float s;
  s = 0.0;
  for (i=...; i<...; i++) {
    s += a[i];
  }
}

```

example calculates
thread-individual sums

useless, from a practical point
of view. But bear with me -
we'll fix this, eventually



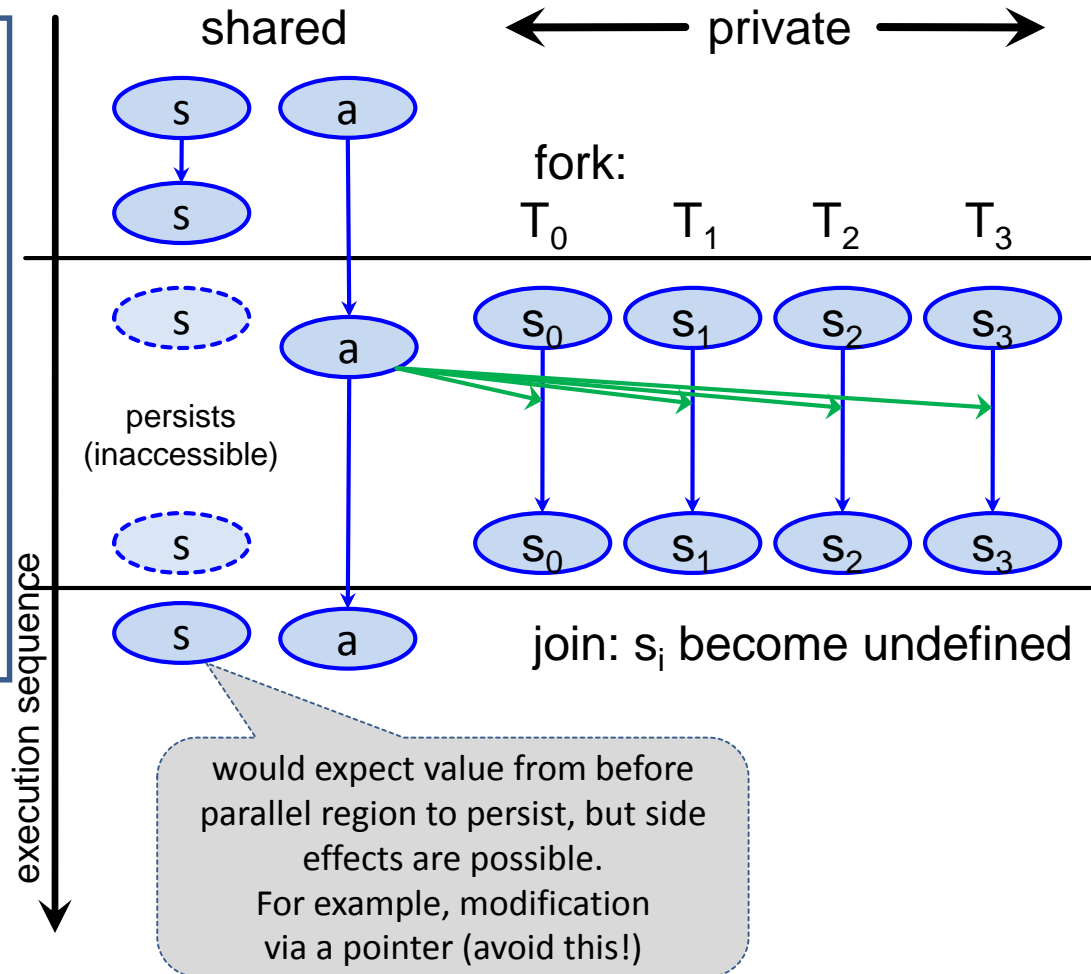
- **Block-local variables in C/C++**
 - are automatically private

Note: I would expect the same behaviour for the Fortran 2008 BLOCK construct, but this is currently not specified in the OpenMP standard

```

Fortran
real :: s
real :: a(:)
integer :: i
s = ...
!$omp parallel private(s) &
!$omp                shared(a)
s = 0.0
do i = ..., ...
  s = s + a(i)
end do
!$omp end parallel
... = ... + s
  
```

- **Masking occurs**
 - for privatized variables declared outside the parallel region
- **Loop variables**
 - are always private



If **s** were shared, the program would have a race condition.

Serial

Fortran

```
DO k = 1, n
  DO j = 1, n
    r(j) = r(j) + a(j, k) * x(k)
  END DO
END DO
```

C

```
for (k=0; k<n; k++) {
  for (j=0; j<n; j++) {
    r[j] = r[j] + a[k*n+j] * x[k];
  }
}
```

OpenMP parallel

```
!$omp parallel
!$omp do
DO j = 1, n
  DO k = 1, n
    r(j) = r(j) + a(j, k) * x(k)
  END DO
END DO
!$omp end do
... = r(...)
!$omp end parallel
```

implicit **barrier**

all threads synchronize

```
#pragma omp parallel
{
  #pragma omp for
  for (j=0; j<n; j++) {
    for (k=0; k<n; k++) {
      r[j] = r[j] + a[k*n+j] * x[k];
    }
  }
  ... = r[...];
}
```

applies to j-loop

no race condition
against previous
definitions

■ Slicing of iteration space

- „loop scheduling“
- default behaviour is implementation dependent
- usually as equal as possible chunks of largest possible size, one chunk per thread

■ In the example,

- slicing is done as shown some slides earlier
- loop order was switched to avoid having many synchronizations

■ Additional clauses

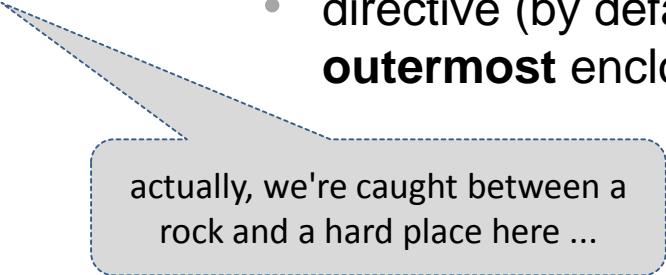
- on both !\$OMP DO and !\$OMP END DO will be discussed in another talk

■ Restrictions on loop structure

- Trip count must be **computable** at entry to loop
- **Disallowed:**
C style loops modifying the loop variable in the loop body, or using a non-evaluable exit condition, or Fortran DO WHILE loop;
- loop body must be a well-formed structured block with single entry and single exit point

■ Note:

- directive (by default) acts only on **outermost** enclosed loop



actually, we're caught between a rock and a hard place here ...

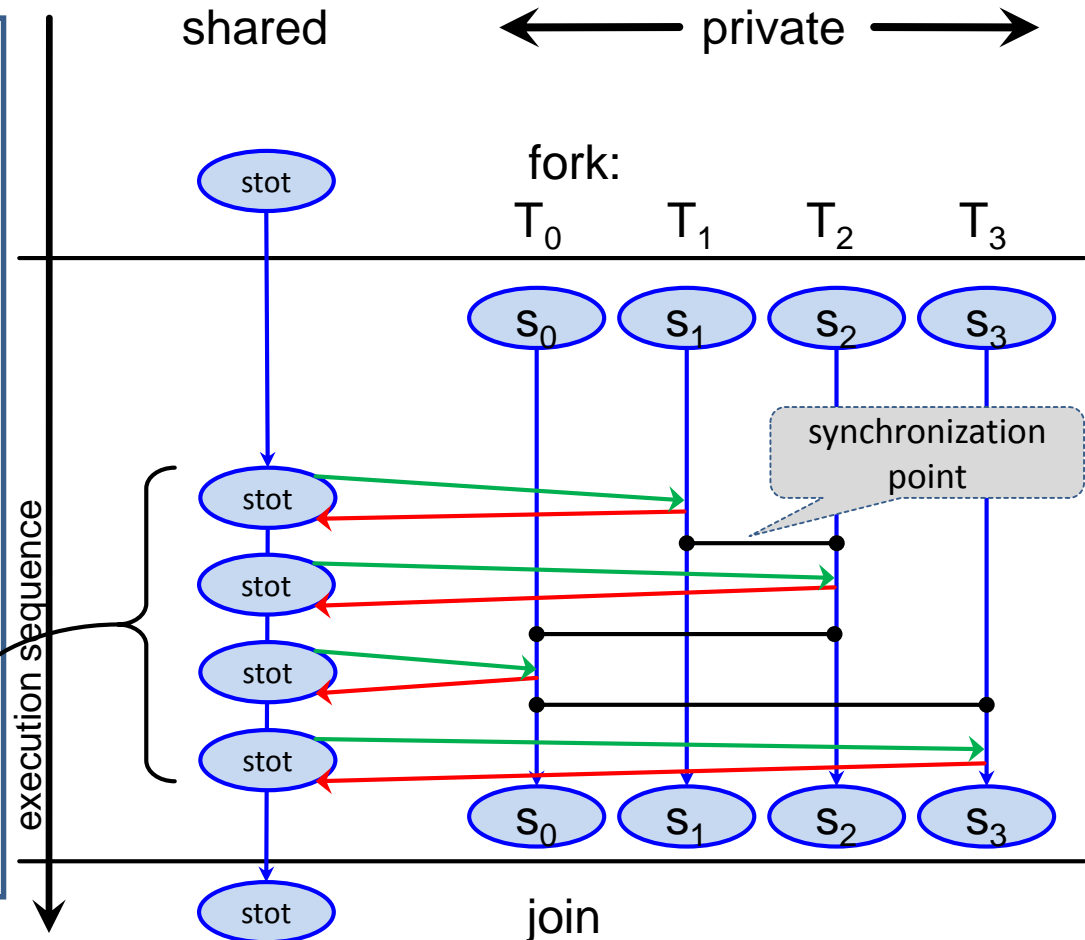
Fortran

```

real :: s, stot
real :: a(:)
integer :: i
  stot = 0.0
!$omp parallel private(s) &
!$omp      shared(a, stot)
  s = 0.0
!$omp do
  do i = 1, size(a)
    s = s + a(i)
  end do
!$omp end do
!$omp critical
  stot = stot + s
!$omp end critical
!$omp end parallel
  
```

updates are now
synchronized

parallel array summation



- Only one thread at a time can execute a **critical region**
 - others must wait → code in region is **effectively serialized**


```
float stot;
stot = 0.0;
#pragma omp parallel \
    shared(a, stot)
{ int i; float s;
  s = 0.0;
#pragma omp for
  for (i=0; i<N; i++) {
    s += a[i];
  }
#pragma omp atomic update
  stot += s;
}
```

parallel array summation

■ Properties of atomic operations

- the **atomic** directive applies only for a **single update** to a **scalar** shared variable of intrinsic type
- this way of updating is safe when executed concurrently
- otherwise, no synchronising effect imposed by semantics
- if hardware atomic instructions are available, likely to be more efficient than a critical region

legacy notation
omp atomic
is also permitted



- **Data accessed by can be shared or private**
 - shared data – one instance of an entity available to all threads (in principle)
 - private data – each per-thread copy only available to thread that owns it
- **Data transfer** transparent to programmer
- **Synchronization**
necessary for accessing shared data from different threads to avoid race conditions
 - implicit barrier
 - explicit directive

```
real :: s
```

Fortran

```

s = ...
!$omp parallel &
!$omp firstprivate(s)

```

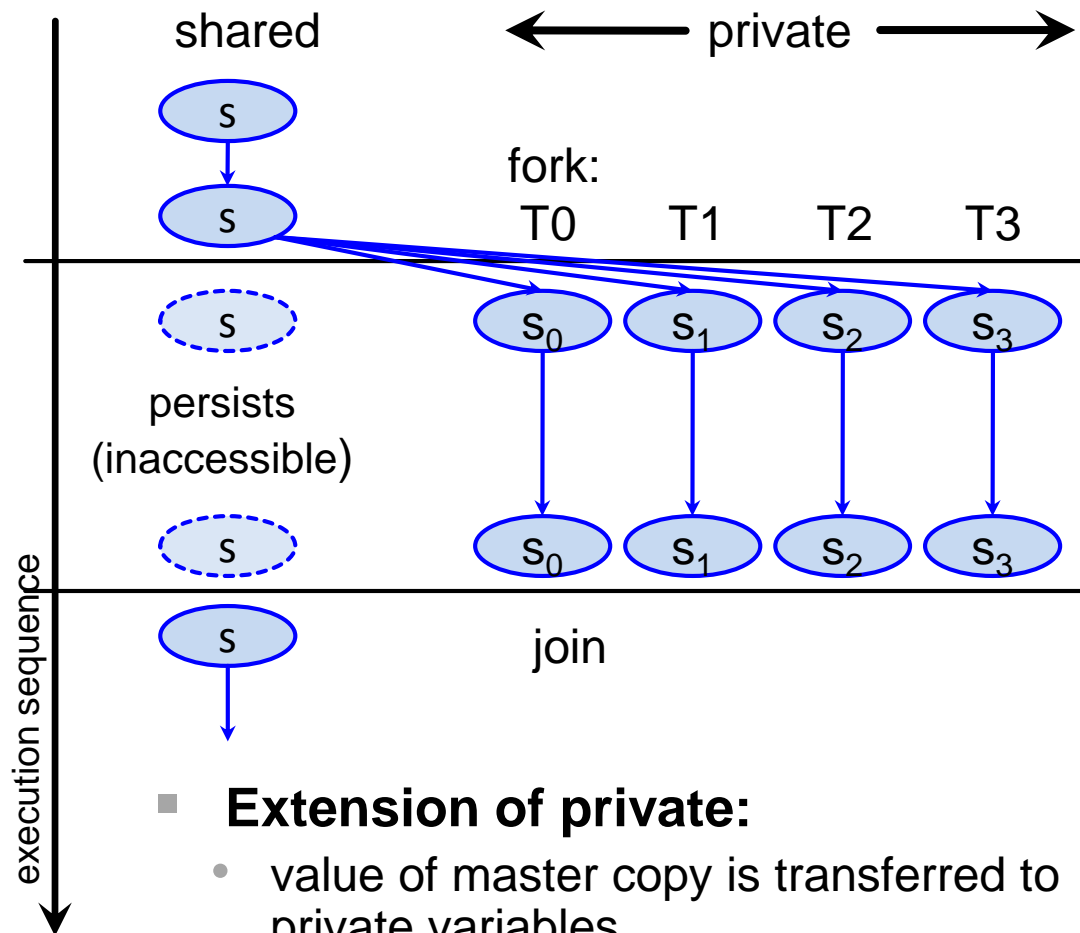
```
... = ... + s
```

uses value from
master copy

```
s = ...
```

```
!$omp end parallel
```

```
... = ... + s
```



Extension of private:

- value of master copy is transferred to private variables
- **restrictions:** not a pointer, not assumed shape, not a subobject, master copy not itself private etc.


```
real :: s
```

Fortran

```

s = ...
!$omp parallel
!$omp do  lastprivate(s)
  do i = 1, n
    s = ...
  end do
!$omp end do
... = ... + s
!$omp end parallel

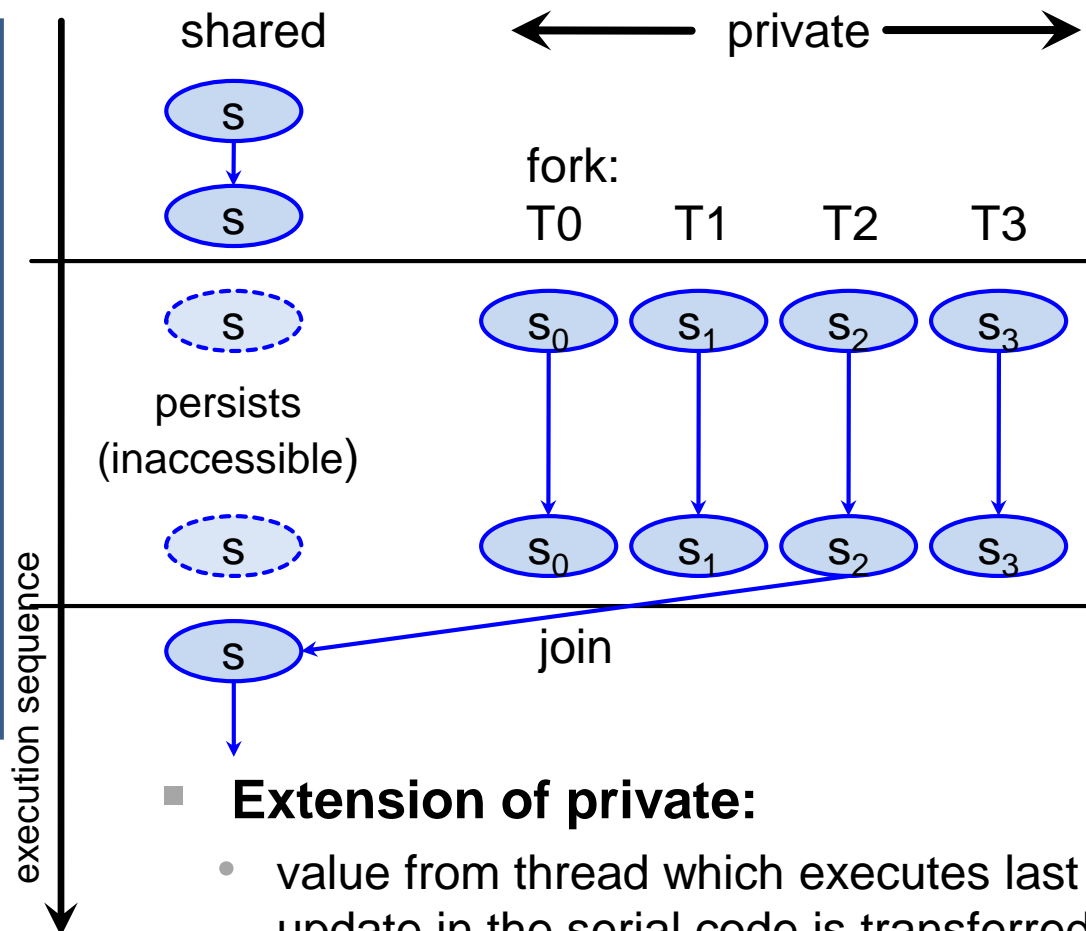
```

on work sharing
directive

s has value produced by
i-loop iteration n

■ When to use?

- as little as possible
- legacy code



■ Extension of private:

- value from thread which executes last update in the serial code is transferred back to master copy
- restrictions similar to **firstprivate**

- **Scoping clauses can be specified for**
 - parallel regions
 - loop work sharing constructs
- **Defaults**
 - apply if no clause is specified
 - may vary by construct, but for the above the following apply:

pre-existing objects are by default **shared**, except for loop variables, which are **private**.

objects declared inside the lexical or dynamic scope of the construct are **private**.

this cannot be changed, of course

- **Recommendation:**
 - specify a **default(none)** clause on each directive that permits scoping:

Fortran

```
!$omp parallel default(none) &  
!$omp  shared(...) private(...) ...  
...
```

other values
are possible

C

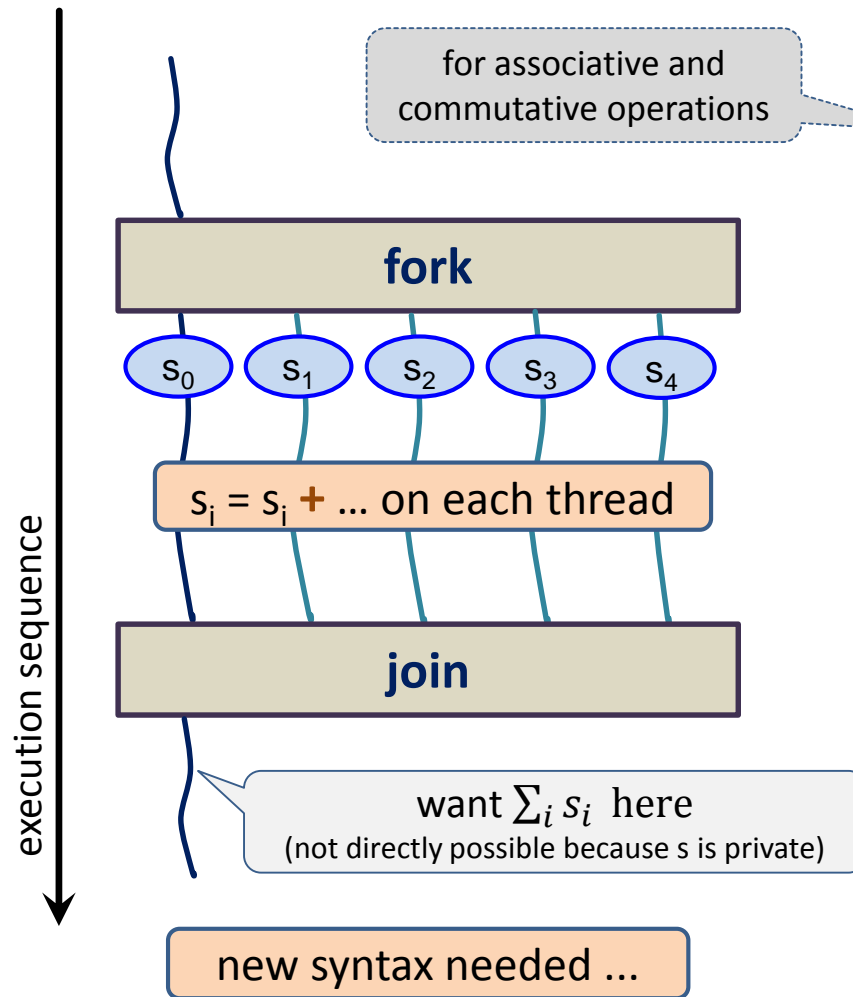
```
#pragma omp parallel default(none) \  
    shared(...) private(...) ...  
...
```

- this **forces** you to explicitly consider and specify scoping for all pre-existing objects

Now: Second exercise session



Reductions




■ **Seen in previous exercise:**

- need for assembling partial results across threads
- up to now: with critical region

■ **OpenMP reductions:**

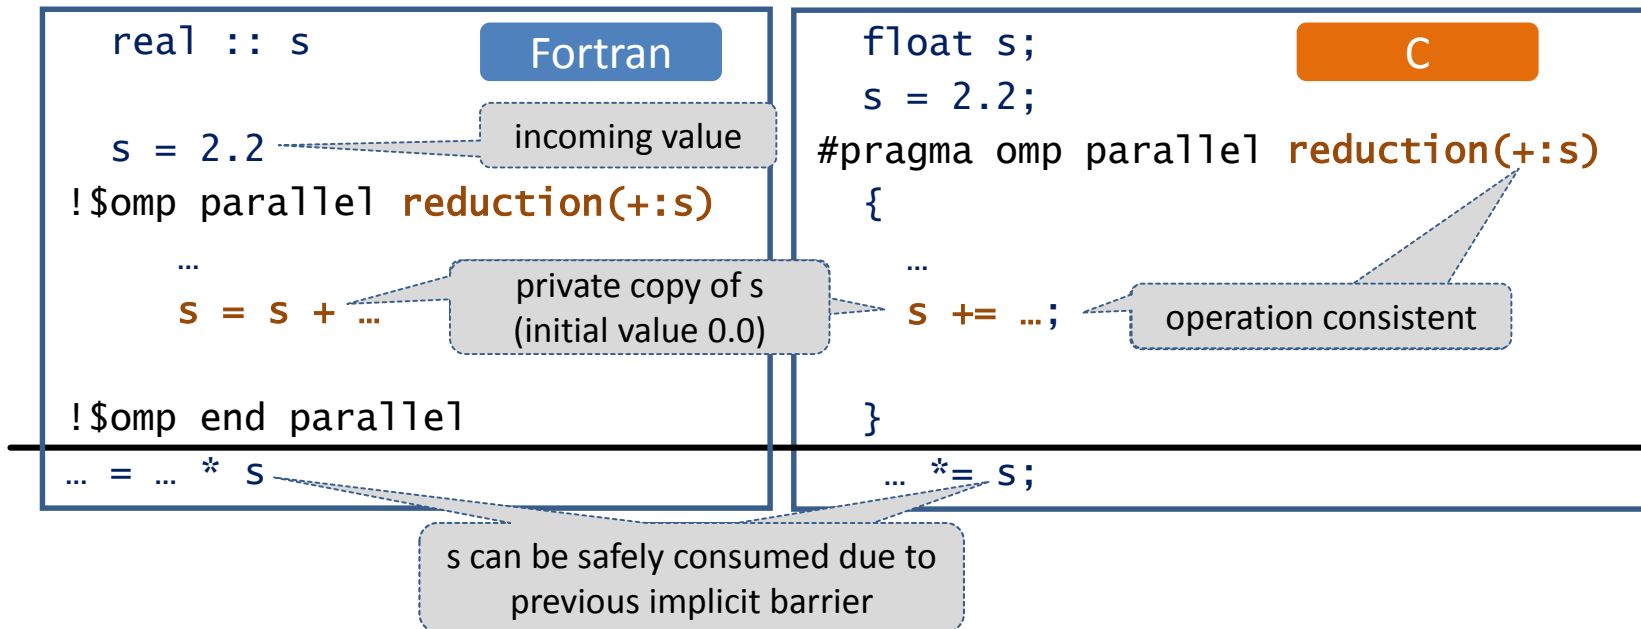
- sometimes more efficient - implementation tunings like



reduce complexity from $O(n_{\text{threads}})$ to $O(\log_2(n_{\text{threads}}))$

- always easier to understand and maintain

■ Example 1: Sum reduction in a parallel region



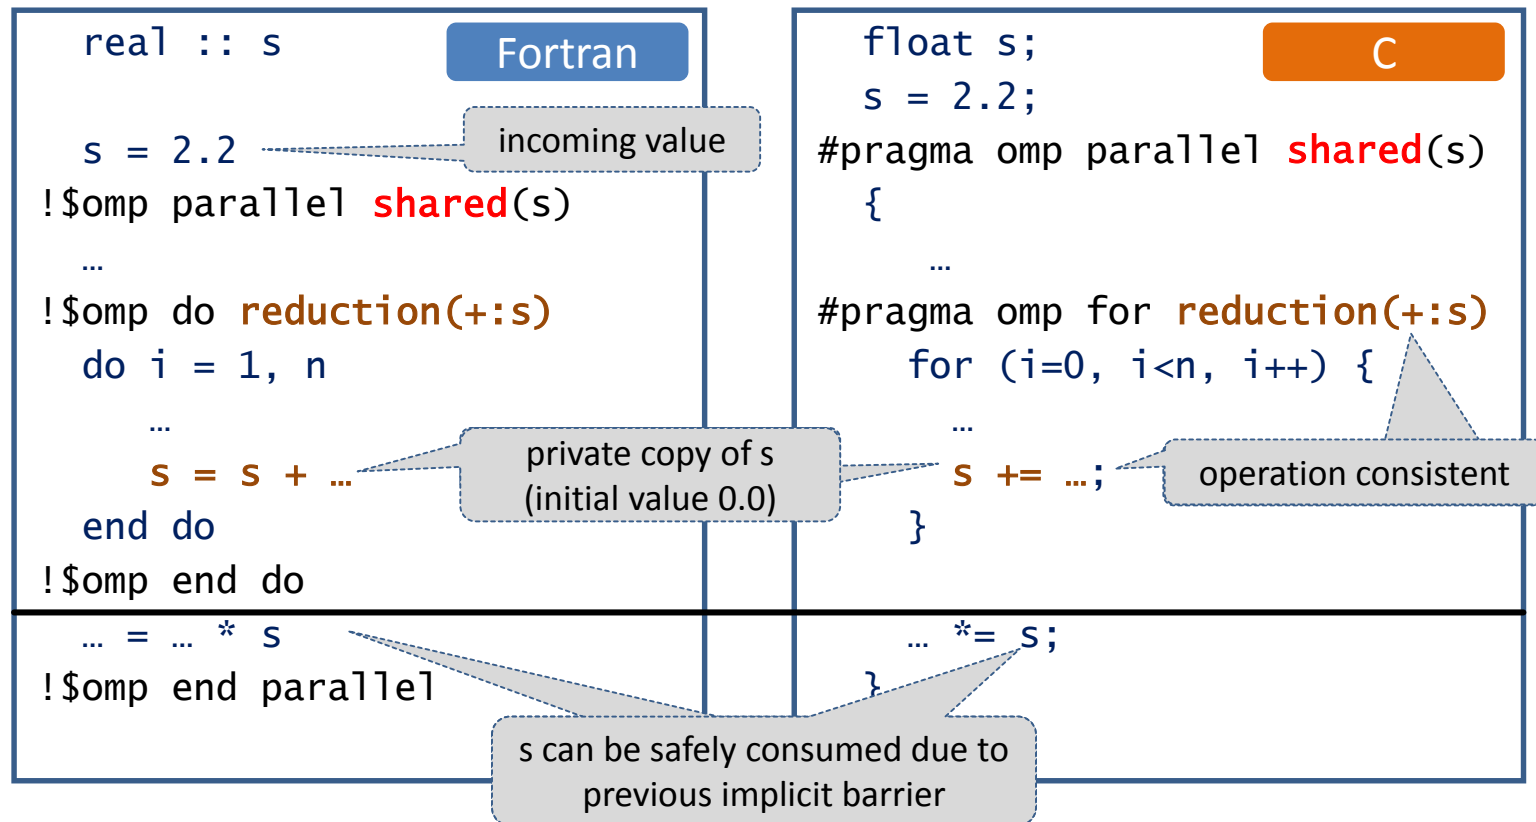
- value of `s` after end of worksharing region: $s_{\text{incoming}} + \sum_i s_i$

■ Note: multiple reductions are permitted

```
!$omp parallel reduction(+:x,y,z)
```

```
!$omp parallel reduction(+:x,y) &
!$omp reduction(*,z)
```

■ Example 2: Sum reduction in a work shared region



- value of s after end of worksharing region: $s_{\text{incoming}} + \sum_i s_i$

- Depends on operation
- Supported intrinsic operations:

Fortran

Operation	Initial value
+	0
-	0
*	1
.and.	.true.
.or.	.false.
.eqv.	.true.
.neqv.	.false.
MAX	-HUGE(X)
MIN	HUGE(X)
IAND	all bits set
IEOR	all bits 0
IOR	all bits 0

C / C++

Operation	Initial value
+	0
-	0
*	1
&	0
	0
^	0
&&	1
	0
MAX	smallest representable value
MIN	largest representable value

```
real :: a(*)  
real :: b(n)
```

Fortran

```
!$omp parallel reduction(+:b) &  
!$omp      reduction(*:a(1:m))  
...
```

must specify
upper bound
(assumed size)

```
float *a;  
float b[N];
```

pointee created
e.g. via malloc()

C/C++

```
#pragma omp parallel \  
reduction(+:b[:]) \  
reduction(*:a[0:m])  
...
```

same as
b[0:N]

■ Example

- reduces complete array b and m elements of array a, elementwise
- uses regular Fortran array section notation

[lower bound : upper bound]

- C example does the same as the Fortran example
- OpenMP-defined sectioning syntax (differs from Fortran):

[lower bound : **length**]

■ General rules:

- array section must be a **contiguous** object (→ no strides permitted)
- dynamic objects must be associated / allocated, and the status must not be modified for the private copies

no deallocate/free within reduction region

■ Using derived types

Fortran

```
type :: fraction
  integer :: numerator, denominator
end type
```

add overloaded operators +, -, * etc.
or even user-defined operators

C

```
typedef struct {
  int numerator, denominator;
} Fraction;
```

provide functions to add, etc.

■ And now we want to write

```
type(fraction) :: af
af = ...
!$omp parallel reduction(+:af)
...
  af = af + ...
!$omp end parallel
```

```
Fraction af;
af = ...;
#pragma omp parallel \
    reduction(+:af)
{
    ...
    Fraction_sum(af, ...);
}
```

- but the compiler will **refuse** to build it („+“ not known to OpenMP) unless further measures are taken ...

```
!$omp declare reduction(+:fraction:omp_out=omp_out+omp_in) &  
!$omp initializer(omp_priv=fraction(0,1))
```

Fortran

```
#pragma omp declare reduction(+:Fraction: \  
Fraction_add(omp_out,omp_in)) \  
initializer(omp_priv=Fraction{0,1})
```

C

■ Combiner

```
declare reduction(<op>:<type>:<combiner>)
```

- connects to operator implementation

Fortran: example defers to overloaded „+“, **C:** references „**Fraction_add**“
special OpenMP parameters **omp_in**, **omp_out** formally describe the two operands for each operation needed

■ Initializer

```
initializer(omp_priv=...) or initializer(function(...))
```

- implements initial value setting for private copies

Fortran: uses (overloaded) structure constructor, **C** similar
special OpenMP parameter **omp_priv** formally describes private copy



More on Work Sharing


Loops and loop scheduling

Collapsing loop nests

Parallel sections

■ Default scheduling:

- implementation dependent
- **typical:** largest possible chunks of as-equal-as-possible size („static scheduling“)



■ User-defined scheduling:

Fortran

```
!$OMP do schedule( static
                   dynamic [,chunk] )
                   guided
```

chunk: always a non-negative integer.
If omitted, has a schedule dependent default value

■ 1. Static scheduling

- `schedule(static,10)`




- minimal overhead (precalculate work assignment)
- default chunk value: see left

■ 2. Dynamic scheduling

- after a thread has completed a chunk, it is assigned a new one, until no chunks are left

`schedule(dynamic, 10)`




- synchronization **overhead**
- default chunk value is **1**


- **Size of chunks in dynamic schedule**
 - too small → large overhead
 - too large → load imbalance
- **Guided scheduling: dynamically vary chunk size.**
 - Size of each chunk is proportional to the number of unassigned iterations divided by the number of threads in the team, decreasing to chunk-size. (default: → 1)
- **Chunk size:**
 - means minimum chunk size (except perhaps final chunk)
 - default value is **1**




- both dynamic and guided scheduling are useful for handling **poorly balanced and unpredictable** workloads.




OMP_SCHEDULE=static



OMP_SCHEDULE=dynamic,10



OMP_SCHEDULE=static,10



OMP_SCHEDULE=guided,10

Decided at run time:

```
Fortran                                     auto
!$OMP do schedule( runtime )
```

- **auto** (automatic scheduling)
 - programmer gives implementation the freedom to use any possible mapping.
- **runtime**
 - schedule is one of the above or the previous two slides
 - determine by either setting OMP_SCHEDULE, and/or calling `omp_set_schedule()` (overrides env. setting)
 - find which is active by calling `omp_get_schedule()`

Examples:

- environment setting:

```
export OMP_SCHEDULE='guided'
export OMP_NUM_THREADS=4
./myprog.exe
```
- call to API routine:

```
omp_set_schedule(omp_sched_dynamic, 4);
#pragma omp parallel
{
  #pragma omp for schedule(runtime)
    for (...) {
        ...
    }
}
```

C

- **Please check your compiler documentation for implementation-dependent aspects**
- **An implementation may add its own scheduling algorithms**
 - code using specific scheduling may be at a disadvantage
 - **recommendation:** Allow changing of schedule during execution
- **If runtime scheduling is chosen and `OMP_SCHEDULE` is not set**
 - execution starts with implementation-defined setting

■ Example: Two nested loops

```

!$OMP do
  do k=1, kmax
    do j=1, jmax
      :
    end do
  end do
!$OMP end do
  
```

Fortran

- assume kmax is 2, and jmax is 3
- then the workshared loop will scale to at most 2 threads

■ Therapy:

- use a collapse clause to improve scaling
- this flattens two (or more) loop nests into a single iteration space

■ Improved example:

```

!$OMP do collapse(2)
  do k=1, kmax
    do j=1, jmax
      :
    end do
  end do
!$OMP end do
  
```

specify nesting level
to collapse

- slicing is performed on the virtual index I_{coll} :

I_{coll}	0	1	2	3	4	5
J	1	2	3	1	2	3
K	1	1	1	2	2	2

sequenced by
serial
execution
order

■ Restrictions:

- rectangular iteration space
- CYCLE/continue in innermost loop only

■ **Example:**




```
!$omp parallel  
!$omp do reduction(+:tsum)  
  do k=1, kmax  
    tsum = tsum + foo(a, b, c)  
  end do  
!$omp end do
```

implicit
barrier

```
...  
... = tsum ...
```

```
!$omp end parallel
```

Fortran

 T_0 performance
slows all others actively executing
 waiting in barrier■ **Assumptions** on code following the synchronization point:

- does not involve **tsum**
- has a load imbalance that is inverse to that of preceding code block

```
!$omp parallel
!$omp do reduction(+:tsum)
  do k=1, kmax
    tsum = tsum + foo(a, b, c)
  end do
!$omp end do nowait
```

no barrier


code not involving tsum

```
!$omp barrier
  ... = tsum ...
```

```
!$omp end parallel
```

Fortran

- **Reduce load imbalance**
 - by removing the barrier via the **nowait** clause
- **Assure code correctness**
 - may require explicit barrier directive before **tsum** (or other modified shared variable) is accessed

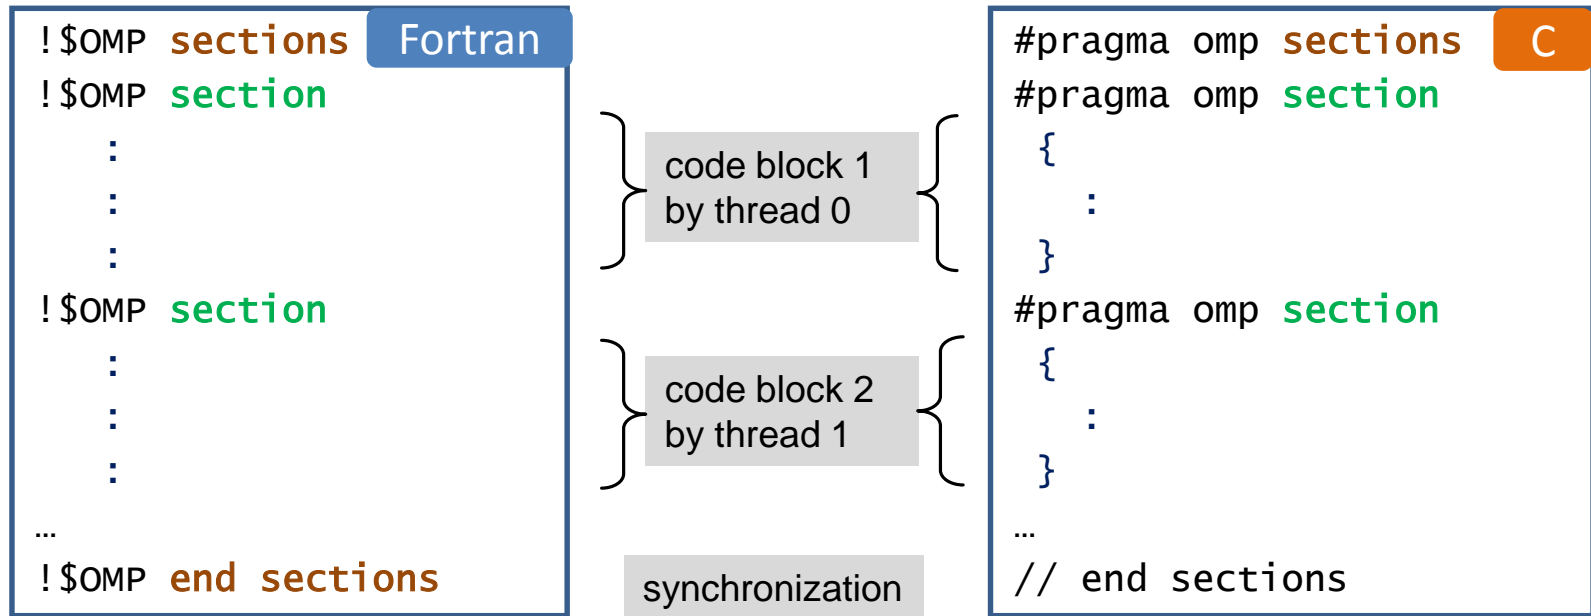


```
#pragma omp for reduction(+:tsum) \
  nowait
{ ... }
```

C

■ Non-iterative work-sharing construct

- distribute a static set of structured blocks



- each block is executed **exactly once** by one of the threads in the team

■ Allowed clauses on sections:

- private, first/lastprivate, reduction, nowait

■ Restrictions:

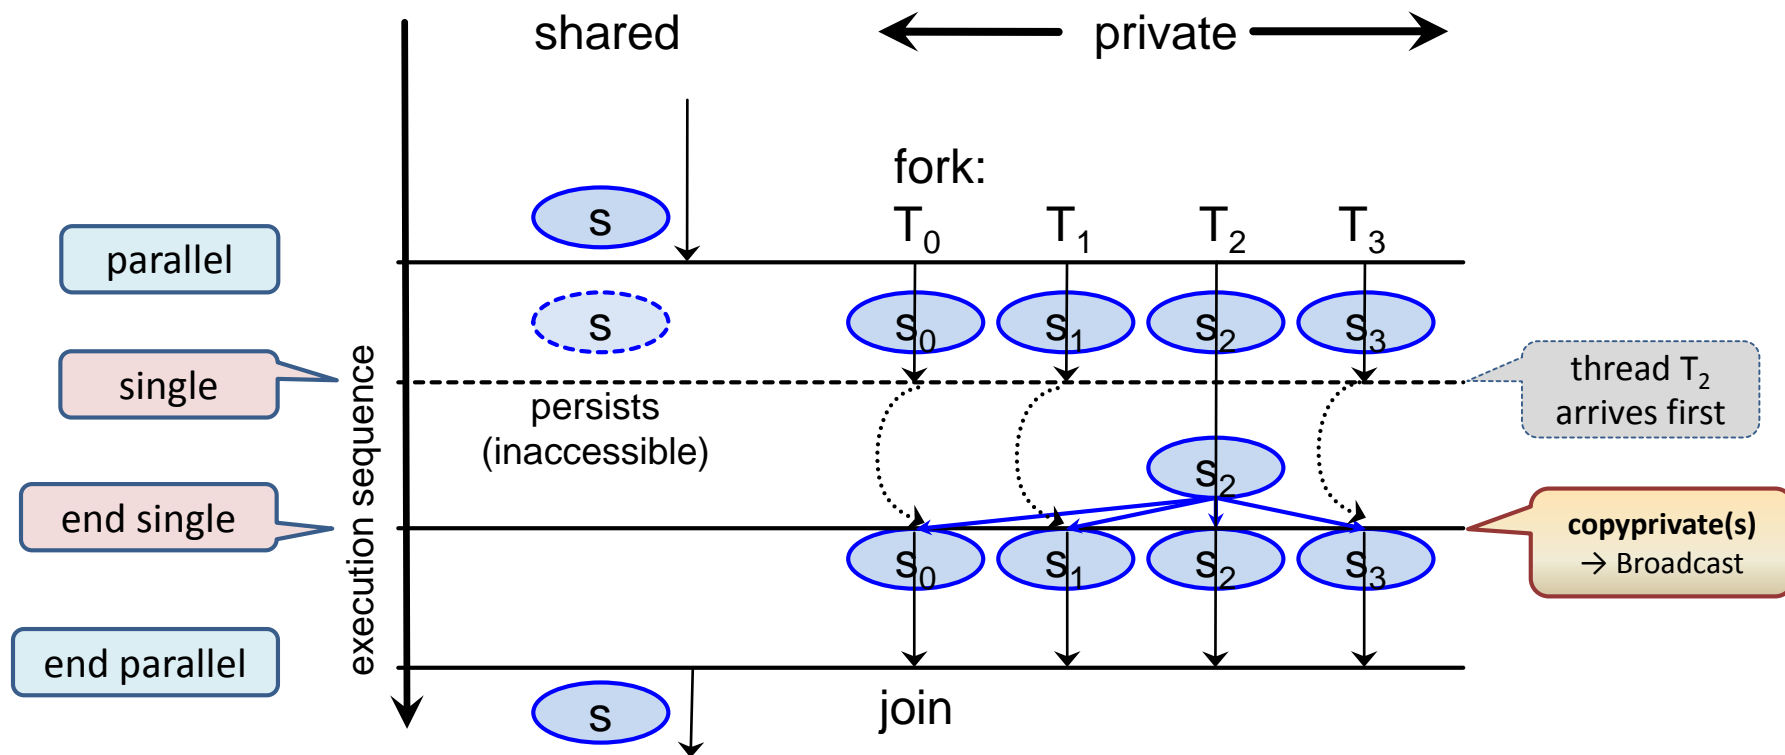
- **section** directive must be within lexical scope of **sections** directive, and directly enclosed (no interleaved language construct is permitted)
- **sections** directive binds to innermost enclosing parallel region
→ only the threads executing the binding parallel region participate in the execution of the section blocks and the implicit barrier (if not eliminated with `nowait`)

■ Scheduling to threads

- implementation-dependent
- if there are more threads than code blocks, excess threads wait at synchronization point

■ In modern OpenMP,

- **tasking** provides a much more flexible and scalable way to implement this and much more general patterns → will be treated tomorrow



■ Execution:

- only one thread of the team executes the statements in the block
- others go to the end of the block

■ Synchronization

- of all threads at end of **single** block

```
real :: s
```

Fortran

```
s = ...  
!$omp parallel private(s)  
  
!$omp single  
...  
s = ...  
!$omp end single &  
!$omp copyprivate(s)  
... = ... + s  
!$omp end parallel
```

block executed by
one thread only

```
float s;
```

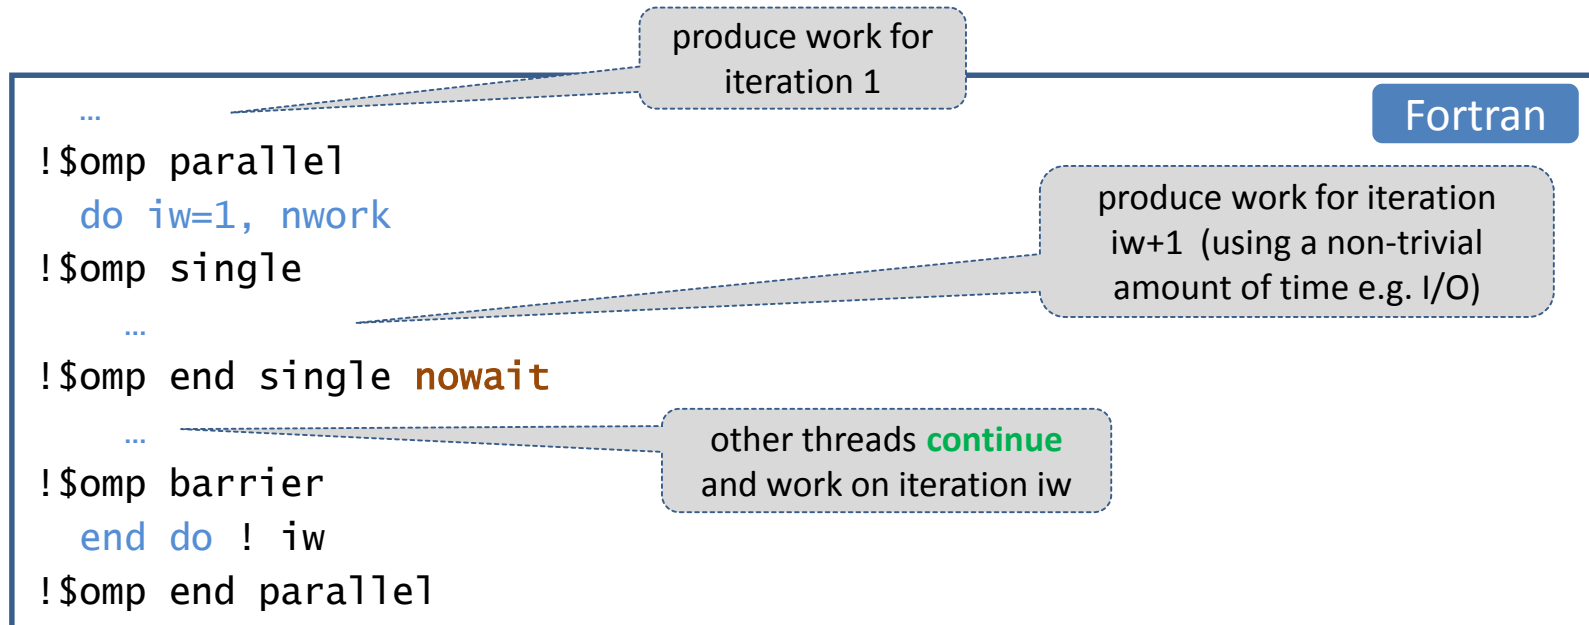
C

```
s = ...;  
#pragma omp parallel private(s)  
{  
#pragma omp single \  
             copyprivate(s)  
{  
    ...;  
    s = ...;  
} // end single  
... = ... + s;  
} // end parallel
```

■ Note:

- update of shared variables inside a single block is safe against subsequent accesses, due to synchronization at the end of that block

- Implement a self-written work scheduler
 - one possible scheme (of many ...), sketched only:



- not the most efficient method
→ preferably use tasking (covered tomorrow); the single construct will be relevant in this context



Global variables and threading

■ Examples:

```
module my_globals
  implicit none
  integer :: my_count
  real, allocatable :: a(:)
  ...
end module
```

Fortran

```
REAL :: A(1000)
INTEGER :: MY_COUNT
COMMON / MY_GLOBS / A, MY_COUNT
```

FORTRAN 77

```
#define NMAX 1000
float a[NMAX];
void my_func() {
  extern float a;
  ...
}
```

C

- Such variables by default have **shared** scope
- The same applies for variables with the **SAVE** (Fortran) or **static** (C) attribute



Implication:

- code using such memory is often **not thread-safe**, unless mutual exclusion is used for accessing the objects

- **When program semantics requires that each thread work on its own copy, privatization is necessary**
 - not exactly the same as private variables → separate syntax needed
- **C:**
 - `#pragma omp threadprivate(list)`
 - list is a comma-separated list of file-scope, namespace-scope, or static block-scope variables that do not have incomplete types
- **Fortran:**
 - `!$omp threadprivate(list)`
 - list is a comma-separated list of named variables and named common blocks. Common block names must appear between slashes.
- **Objects start out with master copy existing only**
 - thread-private copies (with undefined values) spring into existence when the first parallel region is started

■ Copyin clause

- broadcasts object values from master copy to thread-individual copies
- works analogous to the firstprivate clause

```
allocate(a(ndim))  
a(:) = ...  
!$omp parallel copyin(a)  
    ... = a(i) + ...  
    a(i) = ...  
!$omp end parallel
```

Fortran

uses value set on master


■ Subsequent parallel regions:

- thread-individual copies retain their values (by thread) if
 1. second parallel region not nested inside first
 2. same number of threads is used
 3. no dynamic threading is used

Note: none of the potential violations of the above three rules are dealt with in this course

Recommendations:

- Avoid using global variables in the context of threading
- Use object-based design instead



... useful varia

Fortran


```
!$omp master  
    block  
!$omp end master
```

C

```
#pragma omp master  
{ block }
```

- **Only thread zero (from the current team) executes the enclosed code block**
 - there is **no implied barrier** either on entry to, or exit from, the master construct. Other threads continue **without synchronization**
- **Notes:**
 - Not all threads must reach the construct; if the master thread does not reach it, it will not be executed at all
 - this is not a work sharing construct, it only serves for execution control

- **Certain combinations of constructs can be fused**
 - the result is a single construct that behaves as if the two individual ones were tightly nested
 - may be more efficient due to reduced synchronization needs
 - is often easier to read
- **Example: joint "parallel do" (C has "parallel for" here ...)**



- both variants have the same semantics

- Put an "if" clause on a parallel region

```
Fortran
!$omp parallel if (n > 8000)
...
!$omp end parallel
```

process work item of size $O(n^p)$

- specify a scalar logical argument
- may require manual tuning for properly dealing with thread count dependency etc.

- Specific uses:

1. execute serially for small problem sizes (parallel overhead may kill performance)
2. suppress nested parallelism in a library routine:

```
#pragma omp parallel if \
( ! omp_in_parallel() )
{
    ...
}
```

logical / int function from OpenMP run time: are we already parallel in executing scope?

Now: Third exercise session



OpenMP 4.0

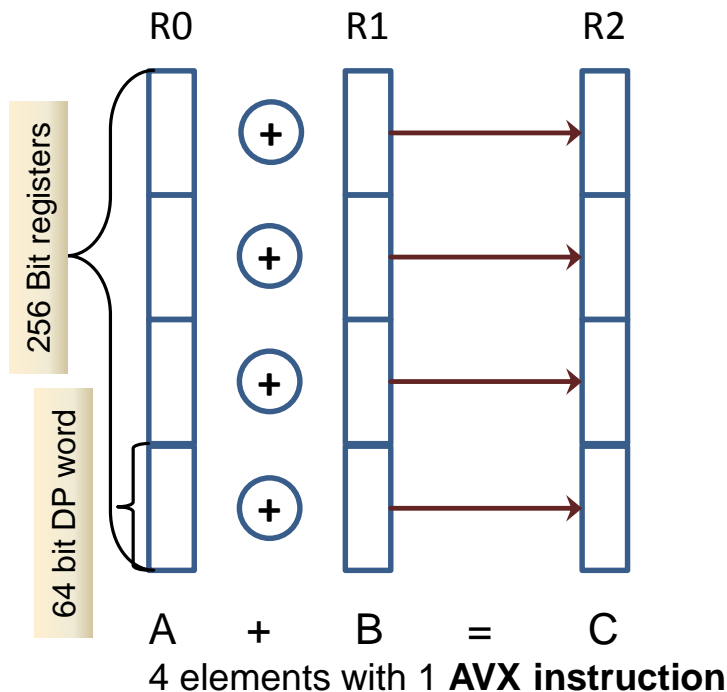
SIMD (vectorization) directives

Optimization of innermost
loop structures

Acknowledgment due to M. Klemm (Intel)

■ Example:

- Sandy Bridge vector unit
- 256 Bit SIMD
- addition of 8 Byte words



■ Instruction capability

- 1 vector add and 1 vector mult per cycle → theoretical Peak 8 Flops/cycle (double precision)

■ LD/ST issue capability for Sandy Bridge

- 4 Words LD/cycle
- 4 Words ST/(2 cycles)
- performance boost depends on algorithm, including its temporal locality properties

■ More recent processors may have more advanced units

- more SIMD lanes
- additional vector operations

- ... programmers had to rely on auto-vectorization,
 - or use **non-portable** extensions
 - programming models (e.g. Intel Cilk Plus)
 - intrinsics (e.g. `_mm_add_pd()`)
 - compiler pragmas

```
#pragma omp parallel for
#pragma vector always
#pragma ivdep
for (int i=0; i<N; i++) {
    a[i] = b[i] + ...;
}
```

C

which may or may not get ignored by the compiler

■ Vectorize a loop nest

- cut into chunks that fit into a SIMD vector register
- without parallelization of the loop body

■ Syntax

```
#pragma omp simd [clause[[,] clause], ...]  
for loops
```

C

```
!$omp simd [clause[[,] clause], ...]  
do loops  
[!$omp end simd]
```

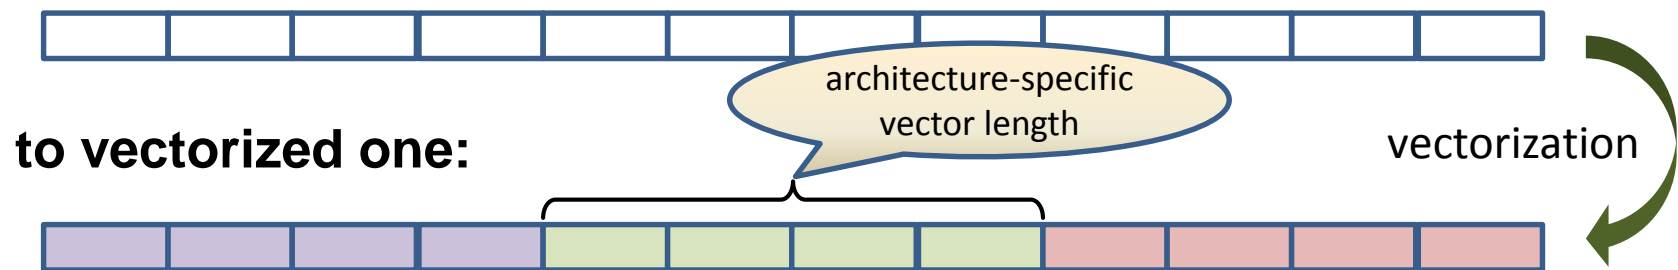
Fortran

- **Scalar product**

```
void sprod(float *a, float *b, int n) {  
    float sum = 0.0f;  
    #pragma omp simd reduction(+:sum)  
    for (int k=0; k<n; k++) {  
        sum += a[k] * b[k];  
    }  
}
```

C


- **Converts serial element-wise execution**



- Existing ones adapted to SIMD-style execution

- required for more complex loop bodies

- `private (var-list)`




create uninitialized vectors for variables in var-list
(loop iteration variables are private by default)

- `lastprivate (var-list)`

copy last iteration value to variable at the end of the construct

- `reduction (op:var-list)`

create private copies for variables in var-list and apply the reduction operation *op* at the end of the construct



■ safelen (length)

- maximum distance between iterations that can run concurrently without breaking any dependencies

```
#pragma omp simd safelen(5)
for (int k=0; k<n; k++) {
    b[k] = a[k] * b[k-j];
}
```

- programmer assures $j > 5$
- compiler can use a vector length of at most 6

■ linear (list[:linear-step])

- produce private copy of a variable that is in linear relationship with the loop iteration variable: $x_i = x_{\text{start}} + (i - i_{\text{start}}) * \text{linear-step}$

- `aligned (list[:alignment])`
 - specifies that variables in the list are aligned, either by the specified integer value of alignment in units of bytes, or in implementation-specific manner
- `collapse(n)`
 - collapse iteration space of a SIMD loop nest

- **Parallelize and vectorize a loop nest**

- distribute iteration space of loops across threads
- subdivide loop chunks to be processed in SIMD registers

- **Syntax**

C


```
#pragma omp for simd [clause[,] clause], ...  
for loops
```

Fortran

```
!$omp do simd [clause[,] clause], ...  
do loops  
[!$omp end do simd]
```

```
void sprod(float *a, float *b, int n) {  
    float sum = 0.0f;  
    #pragma omp for simd reduction(+:sum)  
    for (int k=0; k<n; k++) {  
        sum += a[k] * b[k];  
    }  
}
```

assume invocation by
all threads executing in a
parallel region



- Function call inside SIMD region

```
float min(float a, float b) {  
    return a < b ? a : b;  
}  
  
float distsq(float x, float y) {  
    return (x - y)*(x - y);  
}  
  
void example() {  
    #pragma omp for simd  
    for (i=0; i<N; i++) {  
        d[i] = min(  
            distsq( a[i],b[i] ),c[i] );  
    }  
}
```

may fail if functions
outside file scope

- Therapy: explicitly declare for use in vectorized loops

- C/C++ syntax

```
#pragma omp declare simd  
function def. or decl.
```


- Fortran syntax

```
!$omp declare simd &  
!$omp (proc-name-list)
```

- clauses are also supported
- causes generation of multi-version code by the compiler


- vectorized versions of generated functions are shown

```
#pragma omp declare simd  
float min(float a, float b) {  
    return a < b ? a : b;  
}
```



```
vec8 min_v(vec8 a, vec8 b) {  
    return a < b ? a : b;  
}
```


```
#pragma omp declare simd  
float distsq(float x, float y) {  
    return (x - y)*(x - y);  
}
```



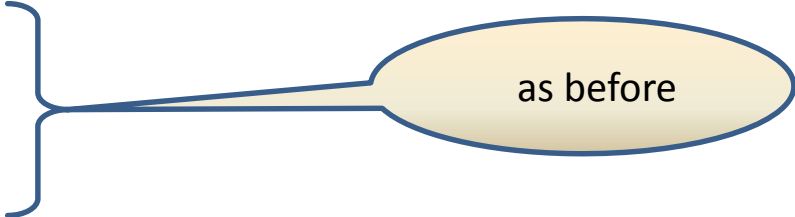
```
vec8 distsq_v(vec8 x, vec8 y) {  
    return (x - y)*(x - y);  
}
```

```
void example() {  
    #pragma omp for simd  
    for (i=0; i<N; i++) {  
        d[i] = min(  
            distsq( a[i],b[i] ),c[i] );  
    }  
}
```

no SIMD directives permitted
inside vectorized functions!



```
vd = min_v(  
    distsq_v (va, vb), vc );
```


- `simdlen (length)`
generate function to support supplied vector length
 - `uniform (argument-list)`
argument has a constant value between iterations of invoking loop
 - `inbranch` vs. `notinbranch`
function always / never called from inside an if statement
 - `linear (list[:linear-step])`
 - `aligned (list[:alignment])`
 - `reduction (op:var-list)`
- 
- as before

- **Case studies on vectorizable applications:**
 - show performance improvements of factor 1.5 – 4.3 compared to auto-vectorized code
 - you may not be as successful, but a 20% performance improvement for 45 min optimization work is also quite nice

- **Resolution of dependencies**
 - may sometimes involve code restructuring and splitting of loops

- **Further features available: combination of device control directives with SIMD**
 - not discussed in this talk

Now: Fourth exercise session



More on Synchronization and Correctness

- Memory model**
- Identifying correctness problems**
- Named critical regions**
- Atomic operations**
- Mutual exclusion with locks**

■ Scenario:

```
real :: a
```

```
a = 0
```

```
!$omp parallel shared(s) num_threads(2)
```

```
a = a + 1
```

```
write(*, '('a on thread ',i0,' is ',i0)') &  
      omp_get_thread_num(), a
```

```
!$omp end parallel
```

```
write(*, '('a after construct is ',i0)') a
```

fix number of threads
for parallel execution

Fortran

possible results
for first write

Thread 0	Thread 1
1	1
2	1
1	2

possible results for second
write: 1 or 2

- the above is **non-conforming**
- data race causes **unpredictable** results to be produced

■ Reason:

- different threads can have different views on same variable: temporary view (in-register value) vs. memory value
- these two views become inconsistent when a thread modifies the variable

■ Flush Operation

- is performed on a set of (shared) variables or on the whole thread-visible data state of a program
- **discards** temporary view:
 - modified values are forced to cache/memory (requires exclusive ownership)
 - next read access must be from cache/memory
- **further** memory operations only allowed after all involved threads complete flush:
 - restrictions on memory instruction reordering (by compiler)

```
!$omp flush [list]
```

recommend to avoid
use of explicit flushes

■ Ensure consistent view of memory:

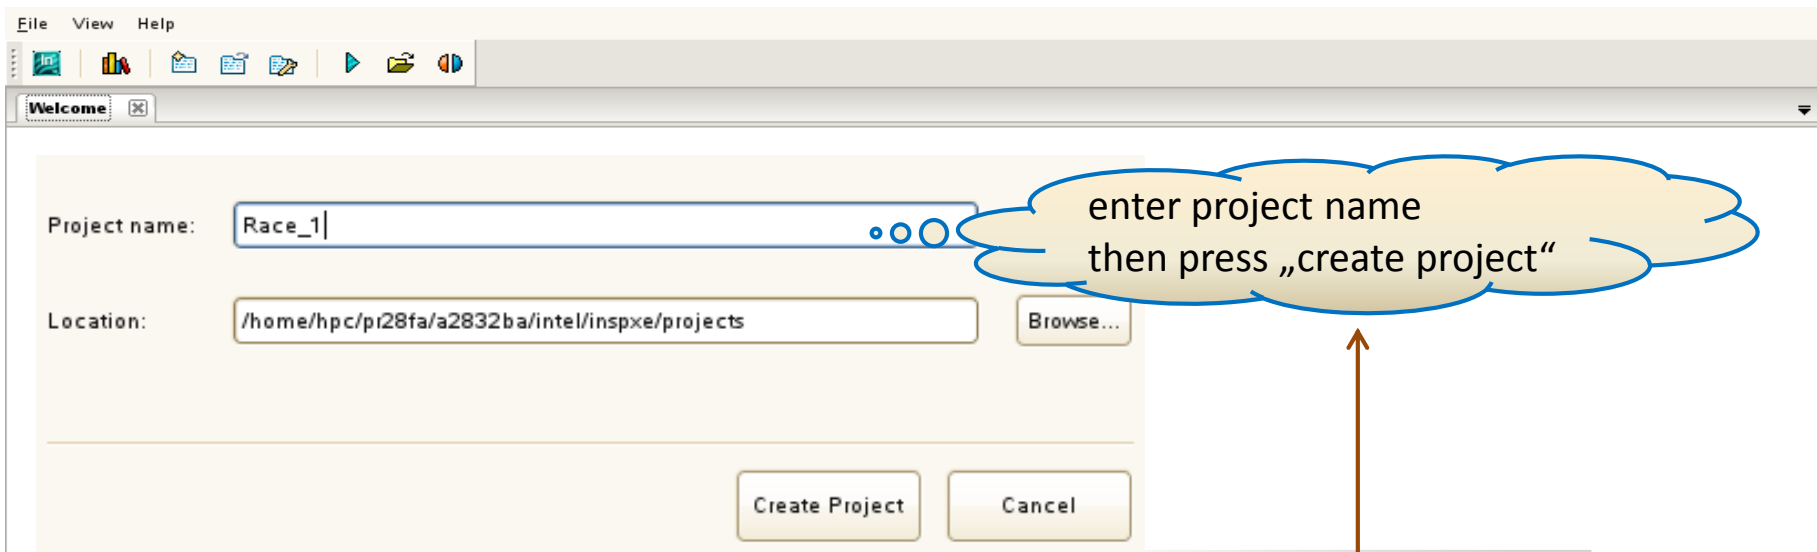
- Assumption: want to write a data item with one thread, read it with another one
 - Order of execution **required**:
 1. thread A writes to shared variable
 2. thread A flushes variable
 3. thread B flushes same variable
 4. thread B reads variable
- The challenge is to assure step 3 happens **after** step 2
 - OpenMP synchronization semantics assure this as well as the necessary flush operations (if correctly used)

- **Example: update via critical region**
 - mutual exclusion is only assured for the statements **inside** the block i.e., subsequent threads executing the block are synchronized against each other
- **If other statements access the shared variable, you may be in trouble:**

```
!$omp parallel shared(x) ...  
:  
!$omp critical  
  x = x + y  
!$omp end critical  
...  
  a = f(x, ...)  
!$omp end parallel
```

Race on read to x.
A barrier is required **before** this statement to assure that all threads have executed their atomic updates

- **OpenMP correctness analysis:**
 - no special compiler option needed (except perhaps `-g`)
 - GUI also for Linux-based system
- **Identify memory issues in addition to threading issues**
 - leaks, dangling pointers etc.
- **Start up GUI**
 - prerequisites: set up environment and possibly stack limit
 - then, invoke the GUI with `inspxe-gui &`
 - command line `inspxe-cl` is also available, but will not be discussed in this talk



File View Help

Project name: Race_1

Location: /home/hpc/pr28fa/a2832ba/intel/inspxe/projects Browse...

Create Project Cancel

enter project name
then press „create project“

Current project: heat_tune1

- ▶ [Threading Error Analysis / Detect Deadlocks and Data Races](#)
- ▶ [Threading Error Analysis / Locate Deadlocks and Data Races](#)
- ▶ [Memory Error Analysis / Detect Leaks](#)
- ▶ [Memory Error Analysis / Locate Memory Problems](#)
- ▶ [New Analysis...](#)

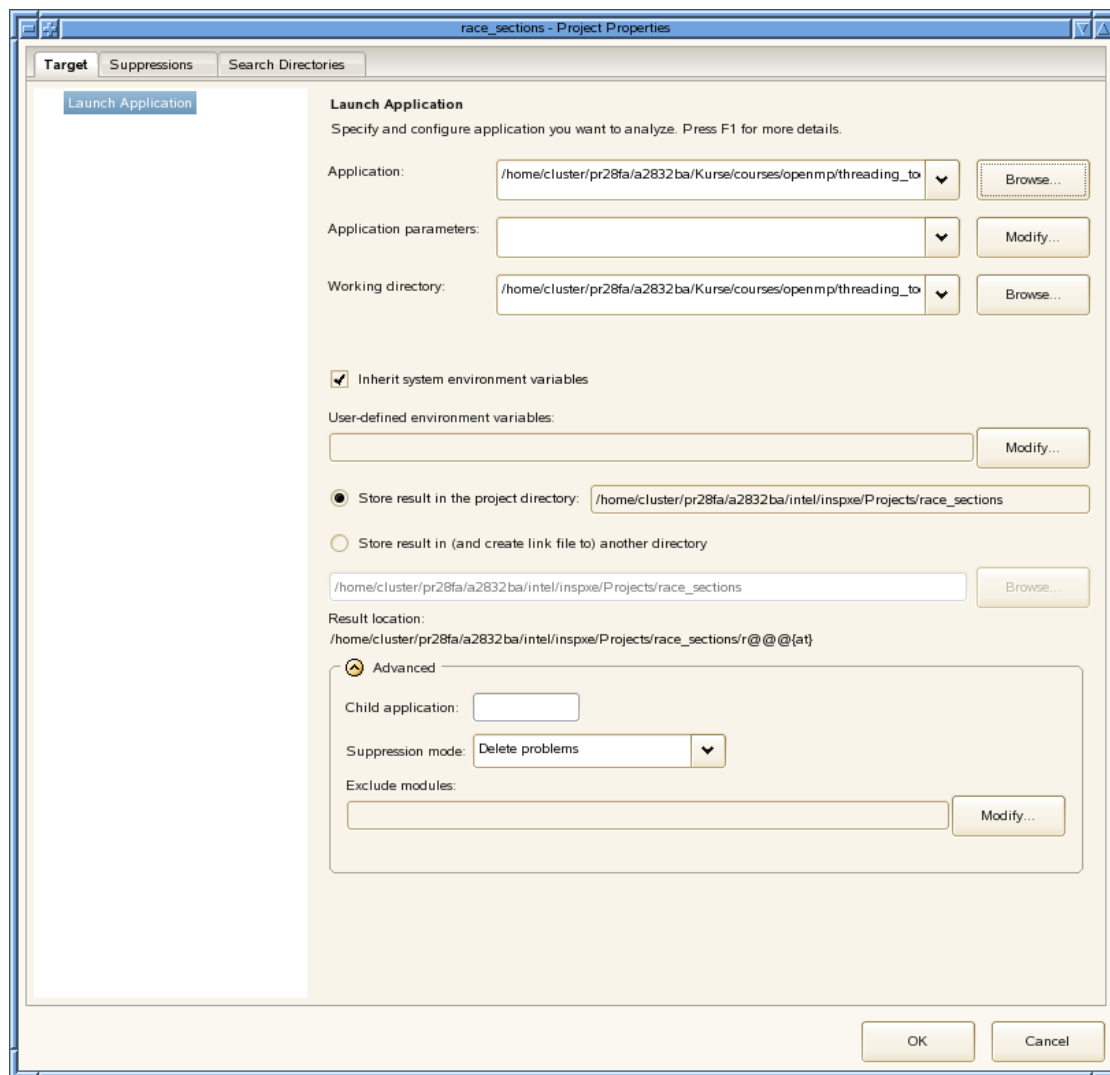
Recent Projects:

- > [tasked_integral_c](#)
- > [tasked_integral](#)
- > [demo_kurs](#)
- > [fp](#)

-  [New Project...](#)
-  [Open Project...](#)

Recent Results:

- > [r004ti3 \[heat_tune1\]](#)
- > [r003ti3 \[heat_tune1\]](#)
- > [r002ti3 \[heat_tune1\]](#)
- > [r001ti3 \[heat_tune1\]](#)
- > [r000ti3 \[heat_tune1\]](#)



The screenshot shows the 'race_sections - Project Properties' dialog box with the 'Launch Application' tab selected. The dialog has three tabs: 'Target', 'Suppressions', and 'Search Directories'. The 'Launch Application' tab contains the following fields and controls:

- Launch Application** section: Specify and configure application you want to analyze. Press F1 for more details.
- Application:** Text field with value '/home/c/cluster/pr28fa/a2832ba/Kurse/courses/openmp/threading_to' and a 'Browse...' button.
- Application parameters:** Text field and a 'Modify...' button.
- Working directory:** Text field with value '/home/c/cluster/pr28fa/a2832ba/Kurse/courses/openmp/threading_to' and a 'Browse...' button.
- ☒ **Inherit system environment variables**
- User-defined environment variables:** Text field and a 'Modify...' button.
- ☒ **Store result in the project directory:** Text field with value '/home/c/cluster/pr28fa/a2832ba/intel/inspxe/Projects/race_sections'
- ☐ **Store result in (and create link file to) another directory**
- Result location:** Text field with value '/home/c/cluster/pr28fa/a2832ba/intel/inspxe/Projects/race_sections/r@@@{at}'
- Advanced** section (collapsed):
 - Child application:** Text field
 - Suppression mode:** Dropdown menu with 'Delete problems' selected
 - Exclude modules:** Text field and a 'Modify...' button

At the bottom of the dialog are 'OK' and 'Cancel' buttons.

- **Needed information:**
 - executable name
(must have been built with OpenMP)
 - executable path
(autocompleted)
 - arguments if needed
by executable
- **Further advanced
settings are possible**

File View Help

Welcome i000ti1 **New Inspector Result**

Configure Analysis Type Intel Inspector XE 2013

Analysis Type

Threading Error Analysis

10x-40x 20x-80x 40x-160x

Detect Deadlocks
Detect Deadlocks and Data Races
Locate Deadlocks and Data Races

Analysis Time Overhead Memory Overhead

Locate Deadlocks and Data Races

Widest scope threading error analysis type. Maximizes the load on the system and the time and resources required to perform analysis; however, detects the widest set of errors and provides context and maximum detail for those errors. Press F1 for more details.

☐ Terminate on deadlock

Stack frame depth: 16

Scope: Normal

☒ Remove duplicates

☒ Analyze without debugger

Run an analysis and report all detected problems. Use to view correctness issues and examine them.

☐ Enable debugger when problem detected

Run an analysis under the debugger and stop every time a problem is detected. Not recommended when running a threading analysis because it can be slow.

☐ Select analysis start location with debugger

Run target application under the debugger with analysis disabled until you choose the application entry point. Set a code breakpoint to stop execution prior to when you want to start analysis. At this point, then use the 'monitor begin-analysis' command to turn on analysis for the duration of the analysis.

Details

Detect deadlocks: Yes
Terminate on deadlock: No
Detect lock hierarchy violations: Yes
Save stack on lock creation: Yes
Cross-thread stack access detection: Hide problems/Show warnings

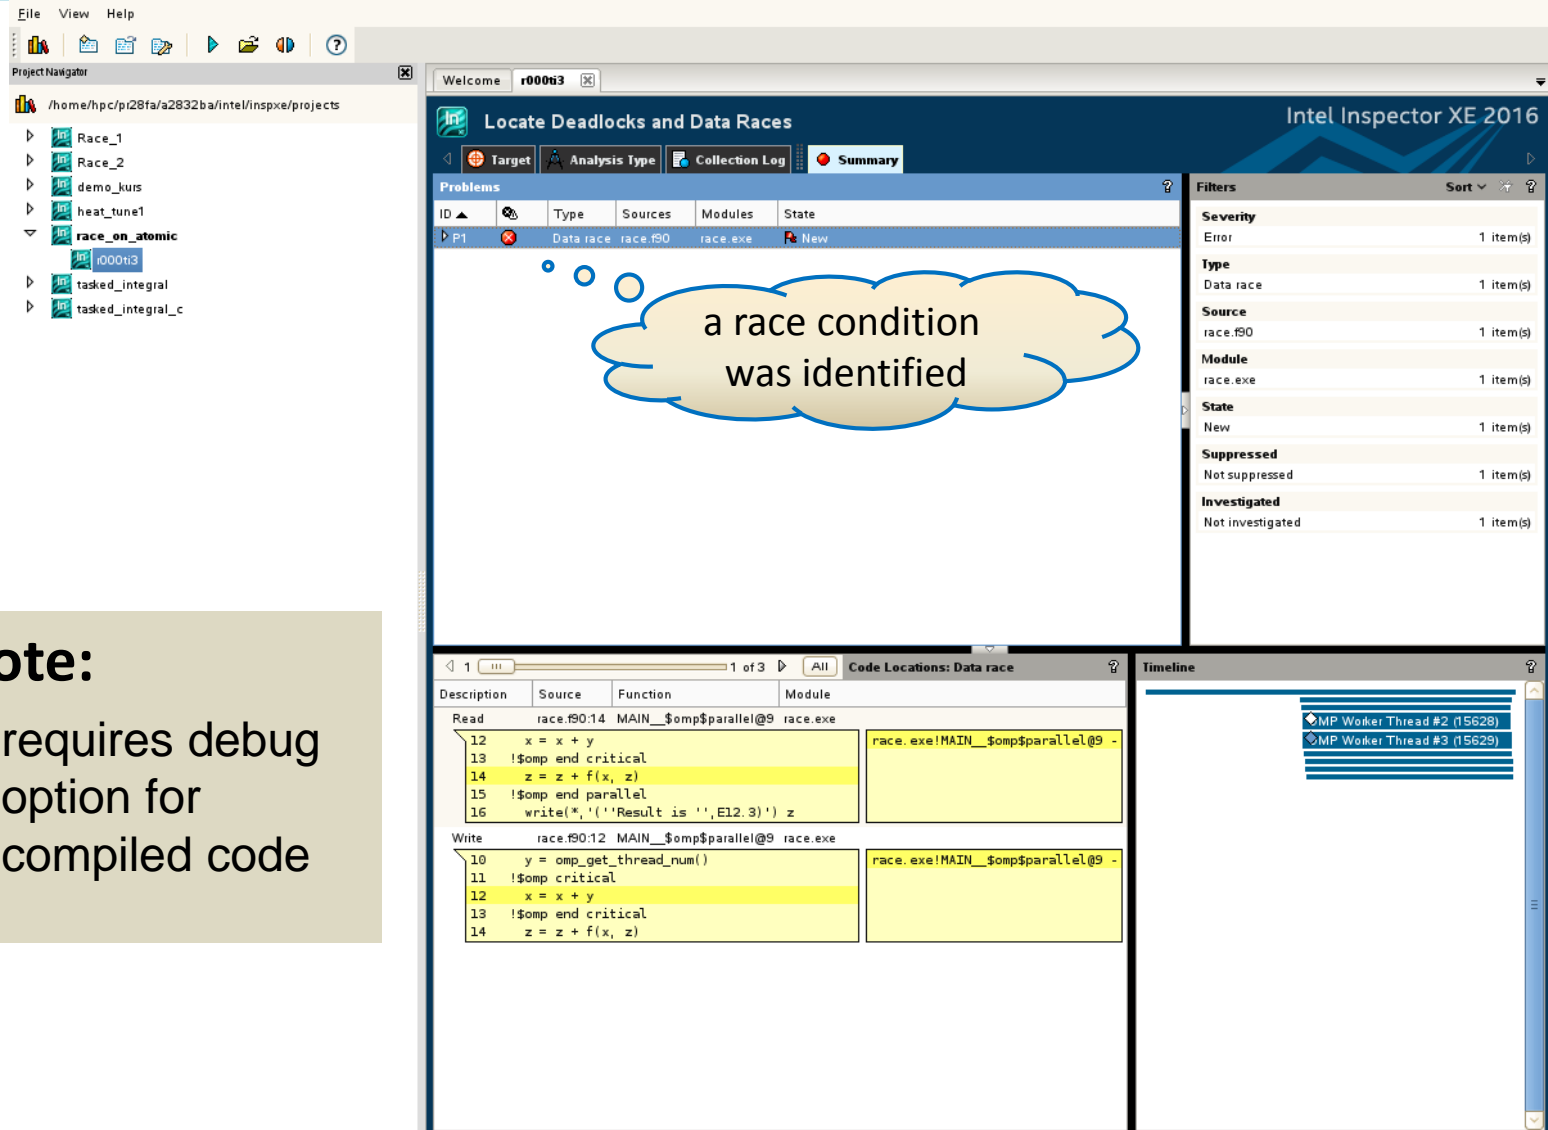
Start
Stop
Reset Leak/Growth Detection
Show Leaks/Growth Now
Close

Project Properties...
Command Line...



Select analysis mode, then start

- here: Threading Error Analysis → locate deadlocks and data races
- note potentially high performance impact



The screenshot shows the Intel Inspector XE 2016 interface. The Project Navigator on the left lists several projects, with 'race_on_atomic' selected. The main window displays the 'Locate Deadlocks and Data Races' analysis results. A table of problems shows a single entry: a 'Data race' in 'race.f90' at 'race.exe', state 'New'. A thought bubble with the text 'a race condition was identified' points to this entry. The right sidebar shows filters for Severity (Error), Type (Data race), Source (race.f90), Module (race.exe), State (New), Suppressed (Not suppressed), and Investigated (Not investigated). The bottom pane shows the 'Code Locations: Data race' view, displaying two code snippets: a 'Read' operation on line 12-16 and a 'Write' operation on line 10-14, both in 'MAIN_\$omp\$parallel@9' of 'race.exe'.

Intel Inspector XE 2016

Locate Deadlocks and Data Races

Target Analysis Type Collection Log Summary

Problems

ID	Type	Sources	Modules	State
P1	Data race	race.f90	race.exe	New

a race condition was identified

Filters

Severity	Count
Error	1 item(s)

Type	Count
Data race	1 item(s)

Source	Count
race.f90	1 item(s)

Module	Count
race.exe	1 item(s)

State	Count
New	1 item(s)

Suppressed	Count
Not suppressed	1 item(s)

Investigated	Count
Not investigated	1 item(s)

Note:

- requires debug option for compiled code

Code Locations: Data race

Description	Source	Function	Module
Read	race.f90:14	MAIN_\$omp\$parallel@9	race.exe
12 x = x + y			
13 !\$omp end critical			
14 z = z + f(x, z)			
15 !\$omp end parallel			
16 write(*, '(Result is ', E12.3)') z			
Write	race.f90:12	MAIN_\$omp\$parallel@9	race.exe
10 y = omp_get_thread_num()			
11 !\$omp critical			
12 x = x + y			
13 !\$omp end critical			
14 z = z + f(x, z)			

Timeline

- MP Worker Thread #2 (15628)
- MP Worker Thread #3 (15629)

The screenshot displays the Intel Inspector XE 2016 interface. On the left, the Project Navigator shows a list of projects, with 'race_on_atomic' selected. The main window is titled 'Data race' and shows a comparison between two threads:

- Read - Thread OMP Worker Thread #3 (15629) (race.exe!MAIN__\$omp\$parallel@9 - race.f90:14)**
- Write - Thread OMP Worker Thread #2 (15628) (race.exe!MAIN__\$omp\$parallel@9 - race.f90:12)**

Both threads are executing the same Fortran code snippet, which includes a parallel region with a critical section. The code is as follows:

```
5 integer :: i
6
7 z = 0.0
8 x = 0.0
9 !$omp parallel private(y) shared(x) reduction(+:z)
10 y = omp_get_thread_num()
11 !$omp critical
12 x = x + y
13 !$omp end critical
14 z = z + f(x, z)
15 !$omp end parallel
16 write(*,('Result is ',E12.3)) z
17 contains
18 pure real function f(a, b)
19 real, intent(in) :: a, b
20 f = a
21 end function
22 end program
```

The interface also shows a Call Stack on the right, indicating the current execution context for each thread.

a) **same** shared variable

thread 0

```
subroutine foo()  
!$omp critical  
  x = x + y  
!$omp end critical
```

thread 1

```
subroutine bar()  
!$omp critical  
  x = x + z  
!$omp end critical
```

Fortran

critical region is **global** → OKb) **different** shared variables

```
subroutine foo()  
!$omp critical  
  x = x + y  
!$omp end critical
```

```
subroutine bar()  
!$omp critical  
  w = w + z  
!$omp end critical
```

Fortran

mutual exclusion not required → unnecessary loss of performance

- **Solution:**
 - use a **named** critical

Fortran

```
subroutine foo()  
!$omp critical (foo_x)  
  x = x + y  
!$omp end critical (foo_x)
```

```
subroutine bar()  
!$omp critical (foo_w)  
  w = w + z  
!$omp end critical (foo_w)
```

mutual exclusion only if same name is used for critical regions acting on different code blocks

- **Note: The atomic directive is bound to the updated variable**
→ problem does not occur when such a directive is used.

- **Assumption:**

- v, w private or shared scalar variables
- x a shared scalar variable

- **Atomic read:**

```
#pragma omp atomic read  
v = x;
```

- **Atomic write:**

```
#pragma omp atomic write  
x = v;
```

- **Atomic capture**

```
!$omp atomic capture  
v = x  
x = x <op> w  
!$omp end atomic
```

- different ordering of statements also allowed

- **Not atomic:**

- evaluation of expressions or updates on v

- **Atomic update:**

- !\$omp atomic update
- same as „traditional“ atomic directive

■ Atomic directives

- permit the programmer to explicitly program with race conditions

■ Rationale for use:

- performance
- tailored synchronizations → will usually require explicit flush operations (not discussed)



Programmer's responsibility

- to assure that no inconsistencies result → must evaluate results from all possible interleavings of execution by different threads
- tools might not be able to observe problems

■ Synchronization effect

- apart from the value change on the variable itself being visible, no synchronization is done
- **sequentially consistent** atomic operations:

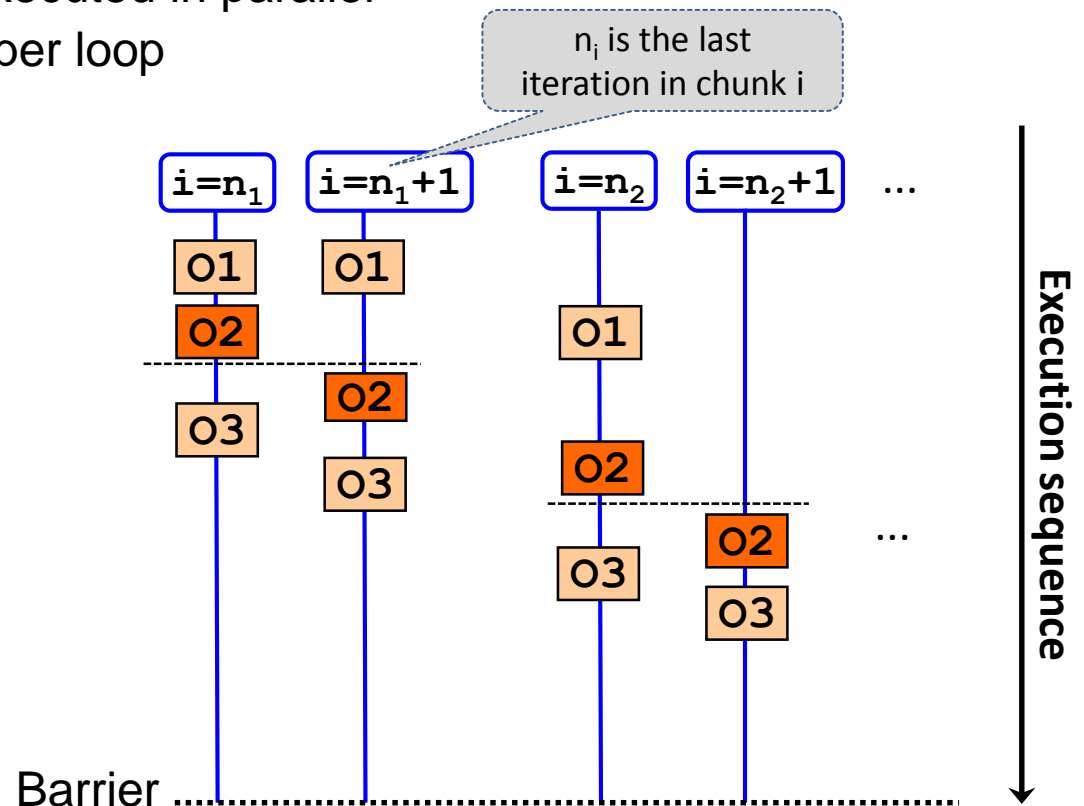
```
#pragma omp atomic \
    seq_cst update
x = x + v;
```

perform a flush on all thread-visible variables (but no synchronization otherwise). Semantics are the same as for such operations in the C++11 standard

- **Statements must be within body of a loop**
 - threads do work **with statements in O2 ordered as in sequential execution**
 - requires `ordered` clause on enclosing loop worksharing directive
 - only effective if code is executed in parallel
 - only **one** ordered region per loop
- **Execution scheme:**

```

!$OMP do ordered
do I=1,N
  o1
!$OMP ordered
  o2
!$OMP end ordered
  o3
end do
!$OMP end do
  
```



■ Loop contains recursion

- dependency requires serialization
- only small part of loop (otherwise performance issue)

```
#pragma omp for ordered
for (i=1;i<n;i+) {
    ... // large block
#pragma omp ordered
    { a(i) = a(i-1)+...; }
} // end loop
```

Fortran

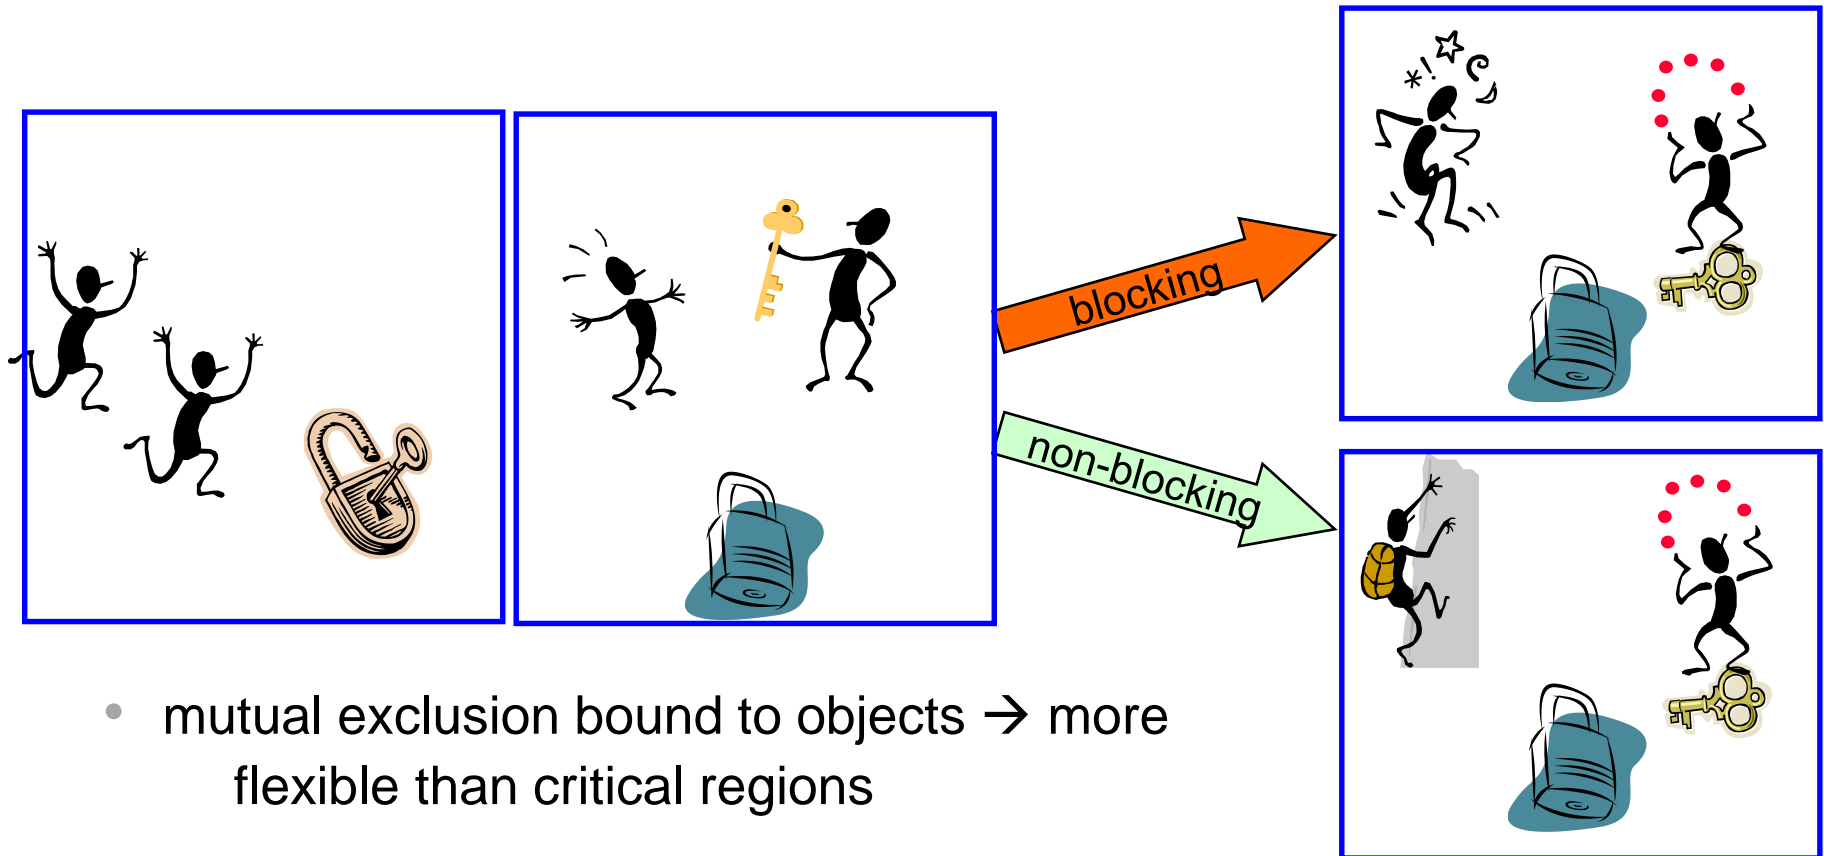
■ Loop contains I/O

- it is desired that content of output (file) be consistent with serial execution

```
!$OMP do ordered
do I=1,N
    ... ! calculate a(:,I)
!$OMP ordered
    write(unit, ...) a(:,I)
!$OMP end ordered
end do
!$OMP end do
```

C

A shared lock variable can be used to implement specifically designed synchronization mechanisms



- **Two variants of locks exist:**
 - simple locks
 - nestable locks (will not be dealt with in detail in this course)
- **Declaration of a lock variable**

```
use omp_lib
```

```
...
```

```
integer(omp_lock_kind) :: a_lock
```

```
integer(omp_nest_lock_kind) :: a_nestable_lock
```

typically an integer capable of
representing an adress

Fortran

```
#include <omp.h>
```

```
...
```

```
omp_lock_t a_lock;
```

```
omp_nest_lock_t a_nestable_lock;
```

C

- **The initial state of a lock variable is "uninitialized"**
 - i.e. it is not actually associated with a lock variable
- **Need to invoke an initialization function on it before it is used**
 - subroutines / void functions provided in OpenMP run time

Name	Purpose
<code>omp_init_lock(omp_lock_t *lock)</code>	initializes an uninitialized lock; the lock variable has the state "unlocked" on return
<code>omp_destroy_lock(omp_lock_t *lock)</code>	destroys a lock that has the state "unlocked".
<code>omp_init_nest_lock (omp_nest_lock_t *lock)</code>	initializes an uninitialized nestable lock; the lock variable has the state "unlocked" on return, and its nesting count is zero.
<code>omp_destroy_nest_lock (omp_nest_lock_t *lock)</code>	destroys a nested lock that has the state "unlocked".

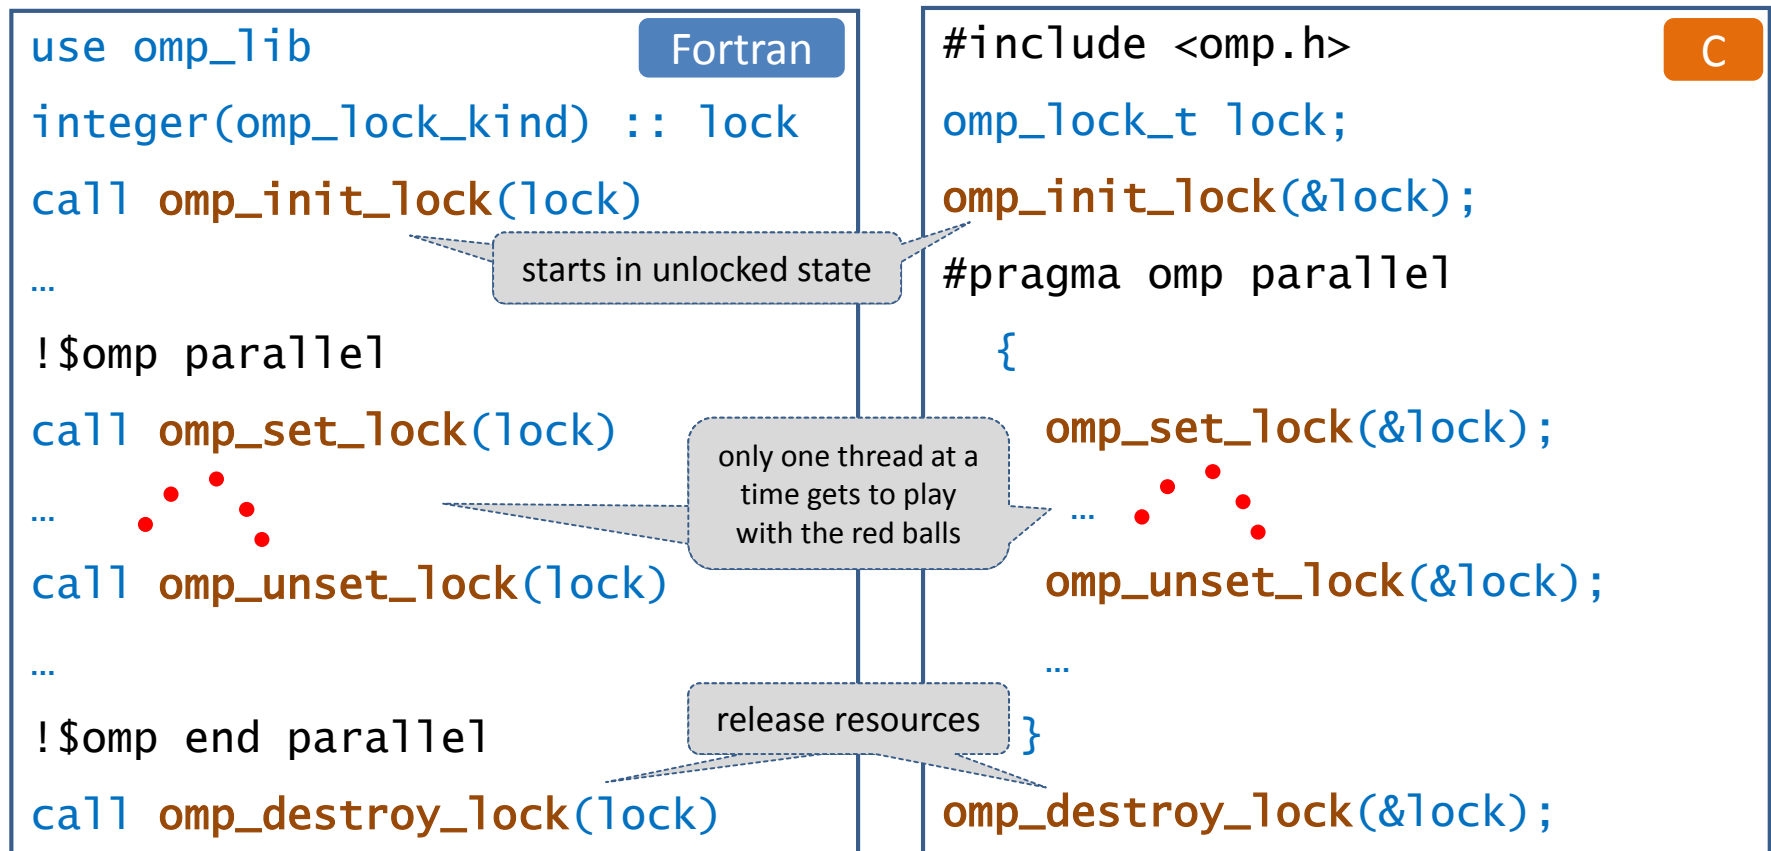
- Fortran: replace `*lock` argument by integer of appropriate kind

- An initialized OpenMP lock can be in one of the states **unlocked**, or **locked**
- The (unique) thread that has successfully acquired the lock is said to **own the lock**
- Only the thread that owns the lock can **release** it, returning it to the unlocked stage.

Name	Purpose
<code>omp_set_lock(omp_lock_t *lock)</code>	If the lock is already locked by another thread, block until the state of the lock changes. If the lock is in the state unlocked, acquire it, setting it to the locked state, and continue execution.
<code>omp_unset_lock(omp_lock_t *lock)</code>	Release the lock that is owned by the executing thread.

- **Notes:**
 - state combinations not described in the table are not permitted (e.g., a thread trying to unset a lock it does not own)
 - the lock variable must be shared in the calling scope

- **Usage pattern analogous to named critical region**
 - programmer is responsible for relationship between lock and objects protected by it



■ Function call signature

```
logical function omp_test_lock(lock)
```

Fortran

```
int omp_test_lock(omp_lock_t *lock)
```

C

- if the lock is already locked by another thread, return "false"
- if the lock has the state unlocked, acquire it (setting the state to locked) and return the value "true".

■ Permits implementing additional concurrency

```
!$omp parallel
```

```
do while (.not. omp_test_lock(lock))
```

```
...
```

```
end do
```

```
...
```

```
call omp_unset_lock(lock)
```

```
!$omp end parallel
```

Fortran

do stuff unrelated
to the red balls

play with the red
balls

```
#pragma omp parallel
```

```
{
```

```
while (! omp_test_lock(&lock)) {
```

```
...
```

```
}
```

```
...
```

```
omp_unset_lock(&lock);
```


```
}
```

C


- **Potential performance issues**
 - locks are a relatively expensive synchronization mechanism
 - lock contention (algorithm dependent)
- **Programming issues**
 - easy to produce deadlock (non-composable against other constructs)
- **Nestable locks**
 - extended semantics for repeated locking (additional nesting count)
- **Locks with hints (OpenMP 4.5)**
 - programmer can specify expected usage pattern, but the actual effect is implementation dependent
 - this is an advanced topic, and success may require special hardware features (transactional processing)

- **Syncbench from the EPCC OpenMP microbenchmarks is used**
 - evaluates the overheads for all synchronizing constructs systematically
 - overhead is what remains even if no workload is processed
- **Showing results as a function of thread count**
 - alternatively, depending on node architecture and used compiler
- **Note order of magnitude**
 - a microsecond typically corresponds to a couple of thousand CPU cycles


Westmere 4-socket node overhead with ICC 15



2-socket results with ICC 15



Westmere 20 thread results



- **Therapy 1:**
 - use the right compiler
 - **note:** x86 does not (yet) support hardware synchronization

- **Therapy 2:**
 - execute serially for small problem sizes
 - conclude parallel execution if not needed any more

- **Therapy 3 (may be most effective):**
 - reduce the synchronization requirements of your algorithm
 - **Examples:** **nowait** clause, or extend parallel regions to reduce number of forks/joins

Now: Fifth exercise session



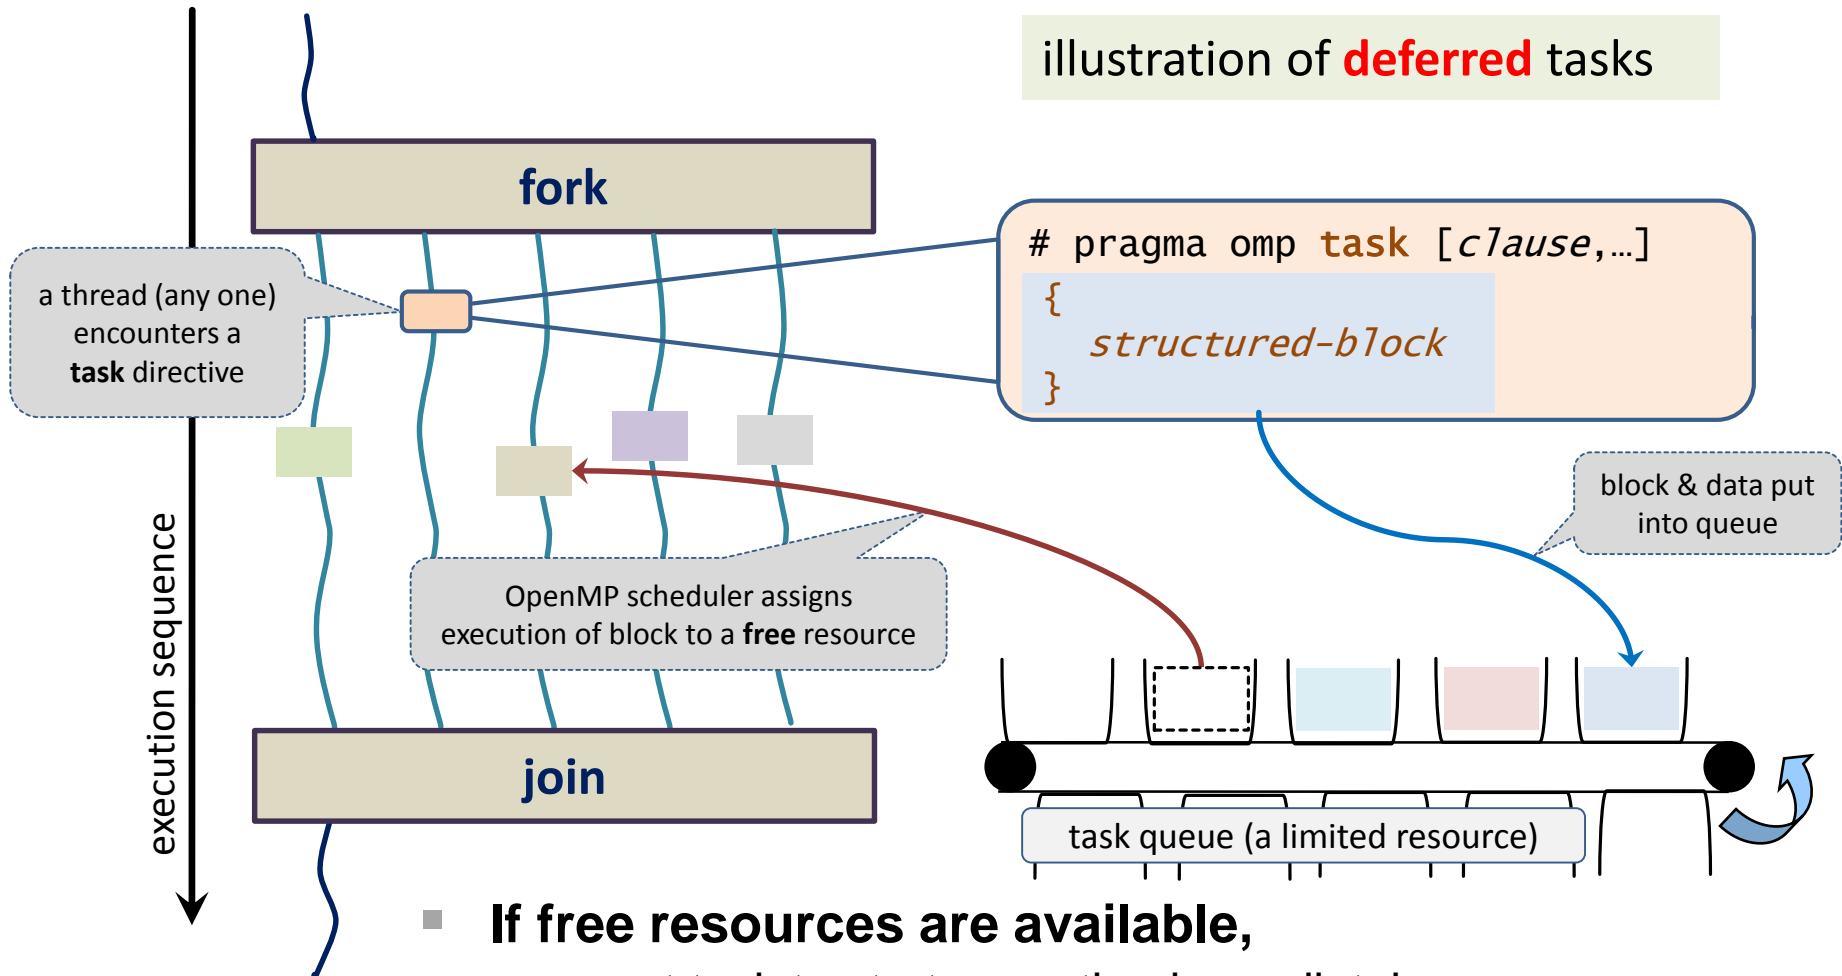
Tasking

**Work sharing for irregular problems,
recursive problems
and information structures**

**Acknowledgement due to L. Meadows/T. Mattson (Intel)
for their SC08 slides**

- **Aim: make OpenMP worksharing more flexible**
- **Semantics:**
 - When a thread encounters a **task construct**, a task is generated from the code of the associated structured block.
 - **Data environment of the task** is created (according to the data-sharing attributes, defaults, ...)
 - The encountering thread may immediately execute the task, or defer its execution.
In the latter case, **any thread in the team** may be assigned the task.
- **Introduced with OpenMP 3.0**
 - additional features and improvements added in later versions of the standard

illustration of **deferred** tasks



- **If free resources are available,**
 - expect task to start execution immediately
- **Task binds to innermost enclosing parallel region**

```
program code_sections
```

Fortran

```
  use mod_functions
```

```
  implicit none
```

```
  real :: a, b
```

```
  integer :: n = ...
```

```
!$omp parallel
```

```
!$omp master
```

```
!$omp task
```

```
  a = function_1(n)
```

```
!$omp end task
```

concurrently executed
if sufficiently many
threads available

```
!$omp task
```

```
  b = function_2(n)
```

```
!$omp end task
```

```
!$omp end master
```

```
!$omp end parallel
```

no synchronization
(different than **single**)

```
  write(*,*) a + b
```

```
end program
```

C

```
int main() {
```

```
  float a, b;
```

```
  int n = ...;
```

```
#pragma omp parallel
```

```
#pragma omp master
```

only thread 0
creates tasks

```
{
```

```
#pragma omp task
```

```
{ a = function_1(n); }
```

```
#pragma omp task
```

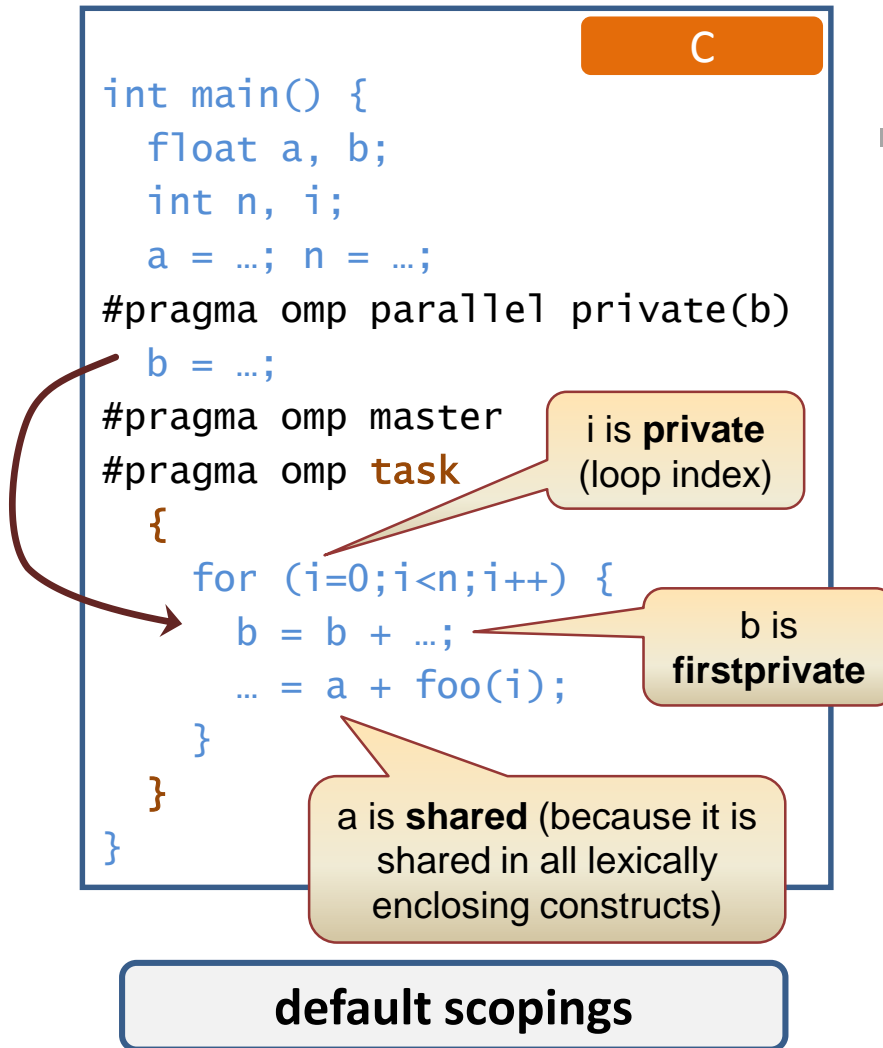
```
{ b = function_2(n); }
```

```
}
```

a and b have
shared scope

```
printf("%f\n", a + b);
```

```
}
```



■ **Recommendation:**

- use a default(none) clause on all task directives
- explicitly specify the scoping for each data object

■ Type declaration


Fortran

```
type :: list
  type(list), pointer :: next => null
  real, allocatable :: data(:)
end type
```

C

```
typedef struct {
  list *next;
  real *data; int n;
} List;
```

■ Data layout



- each list item may carry a different payload
- parallel processing on a per-list-item basis → load imbalance is likely to occur
- the list as a whole is intended to be **shared** (i.e. no copies of payload should be created during processing)


```

subroutine process_list(head)
  type(list), target :: head
  type(list), pointer :: p
!$omp parallel
!$omp single shared(p)
  p => head
  do while (associated(p))
!$omp task firstprivate(p)
    call do_work(p%data)
!$omp end task
    p => p%next;
  end do
!$omp end single nowait
!$omp end parallel
end subroutine

```

Fortran

only one thread
creates taskstask region (includes
procedure execution)synchronization
here → all tasks
done

```

void process_list(list *head) {
  list *p = head;
#pragma omp parallel
{
#pragma omp single \
  nowait shared(p)
  {
    while(p) {
#pragma omp task firstprivate(p)
      { do_work(p->data, p->n); }
      p = p->next;
    }
  } // end single
} // end parallel

```

C

■ Need to have local pointer **p** firstprivate:

- avoid race condition on shared original (vs. subsequent update)
- assure that association status is copied to thread executing the task region

■ When „if“ argument is false,

- the task region is executed immediately by encountering thread (an „undelayed task“)
- but otherwise semantics are the same (data environment, synchronization) as for a „deferred“ task

```
#pragma omp task firstprivate(p) if (sizeof(p->data) > threshold)
    { do_work(p->data); }
```

C

```
!$omp task firstprivate(p) if (size(p%data) > threshold)
    call do_work(p%data)
!$omp end task
```

Fortran

■ User-directed optimization („task pruning“)

- avoid overhead for deferring small task
- avoid creating too many tasks
- cache locality / memory affinity are likely to change

■ Divide and conquer

```
float daq(float *data, int n) {  
    float x1, y1;  
    int n1 = ..., n2 = ...;  
    float *data2 = ...;  
    if (n1 < threshold)  
        { ... }  
    #pragma omp task shared(x1)  
    { x1 = daq(data, n1); }  
    #pragma omp task shared(y1)  
    { y1 = daq(data2, n2); }  
    #pragma omp taskwait  
    return x1 - y1;  
}
```

private at this point

terminate recursion

C

- initial function invocation in a parallel region, usually from a single thread

■ Previous example:


- only sibling tasks are created

■ This example:

- each task creates two child tasks

■ Shared scope for x1, y1:

- start out as private variables
- only newly created tasks share scope with these variables
- shared scope is needed to communicate data outside the task regions



- **The taskwait directive**

- **suspends execution** until immediate child tasks of current task **complete**
(the directive does not apply for descendants of child tasks)

- **Syntax:**

```
!$omp taskwait
```

Fortran

```
#pragma omp taskwait
```

C

- **Needed in example from previous slide**

- avoid race condition of assignments vs. evaluation
- avoid local variables vanishing into thin air while tasks are still executing

■ Possible issues with task scheduling:

- large number of tasks are created → implementation-defined limit on unassigned tasks may be reached
- all currently active tasks reach a synchronization statement → threat of deadlock?

■ Task switching

- permits a thread to suspend a task and start or resume another task **at a task scheduling point**
- for tied tasks, the thread is obliged to resume execution of the suspended task later



Task scheduling points

immediately after generation of a task

at the end of a task region

in implicit or explicit barrier regions

in a taskwait region

in a taskyield region

at the end of a taskgroup region

e.g., a thread that creates lots of tasks may stop doing so and start working on one of them

tasks are tied by default ...

■ Syntax and Semantics

```
!$omp taskyield
```

Fortran

```
#pragma omp taskyield
```

C

- permits (but does not force) task suspension for the current task at the point where the directive is placed

■ Example

- avoid deadlock in a mutual exclusion region
(taken from the OpenMP examples)

```
subroutine foo ( lock, n )  
  use omp_lib  
  integer(kind=omp_lock_kind) :: lock  
  integer :: n  
  integer :: i  
  
  do i = 1, n  
!$omp task  
    call something_useful()  
    do while &  
      ( .not. omp_test_lock(lock) )  
!$omp taskyield  
    end do  
    call something_critical()  
    call omp_unset_lock(lock)  
!$omp end task  
  end do  
  
end subroutine
```

Fortran

■ Example:

Fortran

- a procedure that uses tasking internally without synchronization

```

subroutine p(r,s,...)
  real, intent(inout) :: r, s
  if (...) then
    ...
    r = r + ...
  else
    do while (...)
!$omp task reduction(+:r) shared(s)
      r = r + ...
      call q(s,...)
!$omp end task
    end do
  end if
end subroutine

```

terminate recursion

procedure q might itself do tasking

■ Top-level invocation:

```

real :: ra, sa
ra = 0.0
!$omp parallel shared(ra, sa)
!$omp taskgroup
!$omp master
  sa = ...
  call p(ra, sa, ...)
!$omp end master
  ...
!$omp end taskgroup
  ... = ra * sa + ...
!$omp end parallel

```

assure all created tasks have completed before **ra** is referenced

■ Synchronization

- includes all tasks created inside the region (including descendants)

- **Default behaviour:**

- a task assigned to a thread must be (eventually) completed by that thread → task is **tied** to the thread

- **Change this via the untied clause**

- execution of task block may change to **another** thread of the team at any task scheduling point

```
# pragma omp task untied  
structured-block
```



- **Deployment of untied tasks**

- **Starvation scenario:**

Task switching has caused the task-generating thread to run a long calculation, with the result that all generated tasks were consumed and most threads idle.

If the task that generates the work is untied, a different thread can take over the task-generating workload.



Thread-related semantics used in the untied task region are likely to trip you up, for example ...

- relying on results delivered by `omp_get_thread_num()`
→ may become inconsistent after thread switch
- referencing and defining values stored in threadprivate global variables
→ may access a different copy after thread switch
- using locks or critical regions
→ may, after thread switch, attempt to unlock from a thread which has not taken the lock, etc. Crash may be imminent ...

- **Use of threadprivate data**

- value of threadprivate variables cannot be assumed to be unchanged across a task scheduling point. Can be modified by another task executed by the same thread.

- **Tasks and locks:**

- if a lock is held **across** a task scheduling point, interleaved code trying to acquire it (with the same thread) may cause **deadlock**

- **Tasks and critical regions:**

- similar issue if suspension of a task happens inside a critical region and the same thread tries to access the same critical region in another scheduled task

- **Tools?**

- correctness tools will currently only find some of the issues that can arise

■ Final tasks

- use a **final** clause with a condition on a task directive
- the resulting task is always **undelayed**
- reduces the overhead of placing tasks in the “task pool”
- all tasks created inside task region are also final (different from an **if** clause)
- inside a task block, `omp_in_final()` can be used to check whether the task is final

■ Merged tasks

- using a **mergeable** clause may create a merged task if it is undelayed or final
- a merged task has the same data environment as its creating task region

■ Final and/or mergeable

- can be used for optimization purposes
- e.g. to optimize wind-down phase of a recursive algorithm

Now: Sixth exercise session



Performance:


Architectural aspects

- **What can be expected from the processor architecture?**
 - want at least an estimate for performance limits → avoid „stumbling in the dark“
 - much more detailed node performance engineering and modeling: course by G. Hager and G. Wellein – see <http://moodle.rrze.uni-erlangen.de/moodle/course/view.php?id=300&username=guest&password=guest&lang=en> and references cited within

- **How to exploit the architecture as best as possible**
 - use optimal data access patterns
 - minimize synchronization overhead
 - Account for interactions of OpenMP features with „serial“ optimization techniques (might be compiler optimization or lack thereof!)

■ Performance Characteristics

- determined by memory hierarchy



■ Impact on Application performance: depends on where data are located

- **temporal locality:** reuse of data stored in cache allows higher performance
- **no temporal locality:** reloading data from memory (or high level cache) reduces performance

■ For multi-core CPUs,

- available bandwidth may need to be shared between multiple cores

→ shared caches and memory

■ A small but fast memory area

- used for storing a (small) memory working set for efficient access

■ Reasons:

- physical and economic limitations

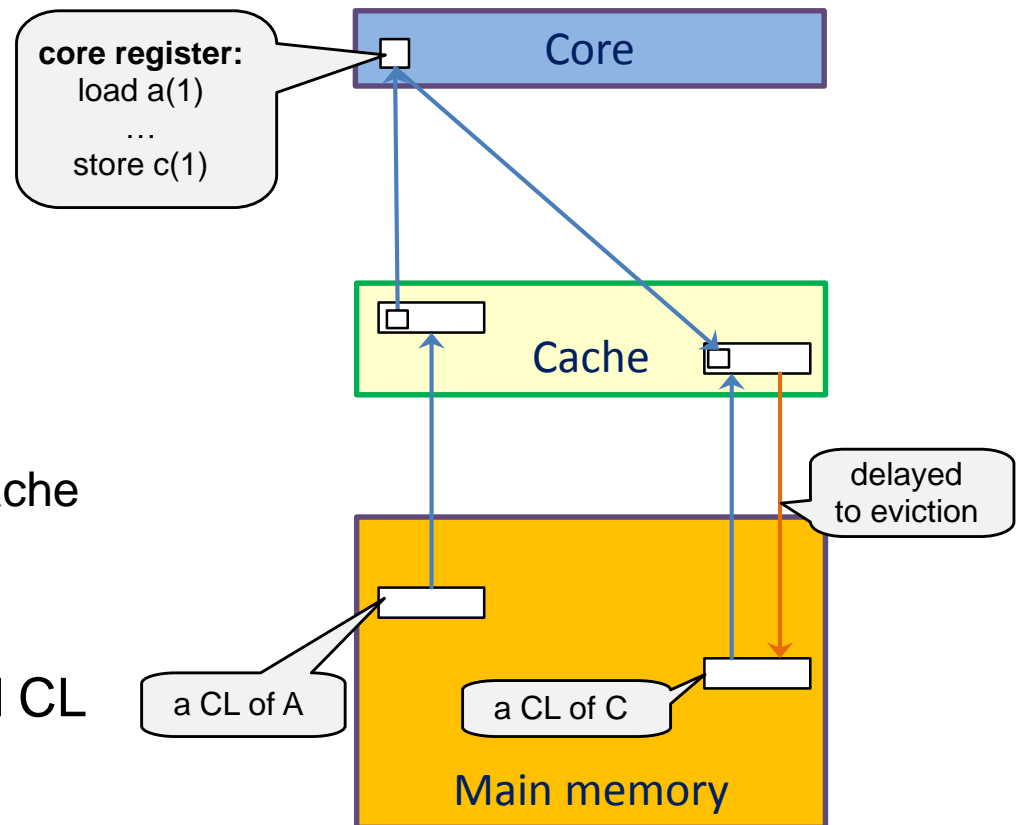
■ Loads (stores) to (from) core registers


- may trigger cache miss → transfer of memory block („cache line“, CL) from memory

■ Cache fills up ...

- usually least recently used CL is evicted

■ Example: $c(:) = a(:) + \dots$





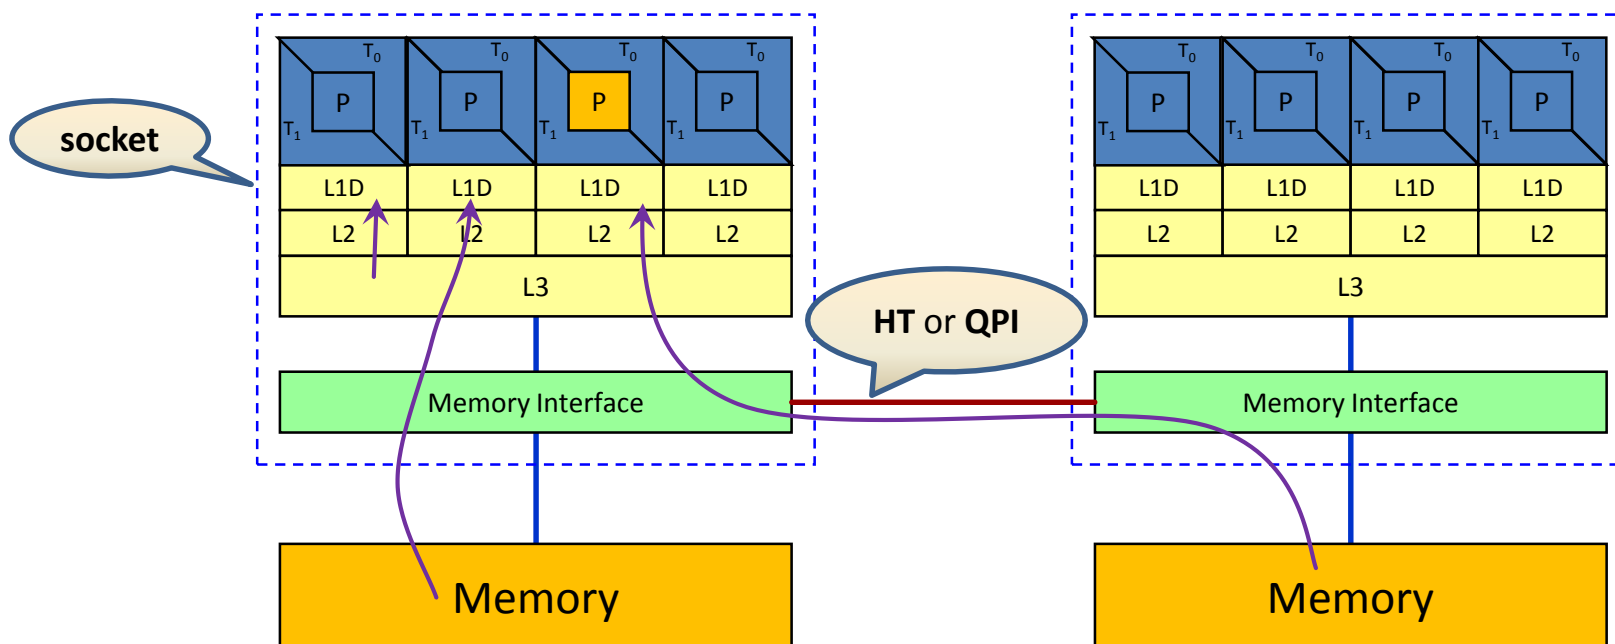
Control of Affinity

NUMA effects

False Sharing

- **multi-core** **multi-threaded** processors with a deep cache hierarchy
- typically, two **sockets** per node

Illustration shows 4 cores per socket. Current sockets have 8 – 14 cores



ccNUMA architecture: „cache-coherent **non-uniform** memory access“

- An implementation might support this:

```
#include <stdio.h>
int main() {
  #pragma omp parallel
  { ...
    #pragma omp parallel
    {
      ...
    }
  }
  return 0;
}
```

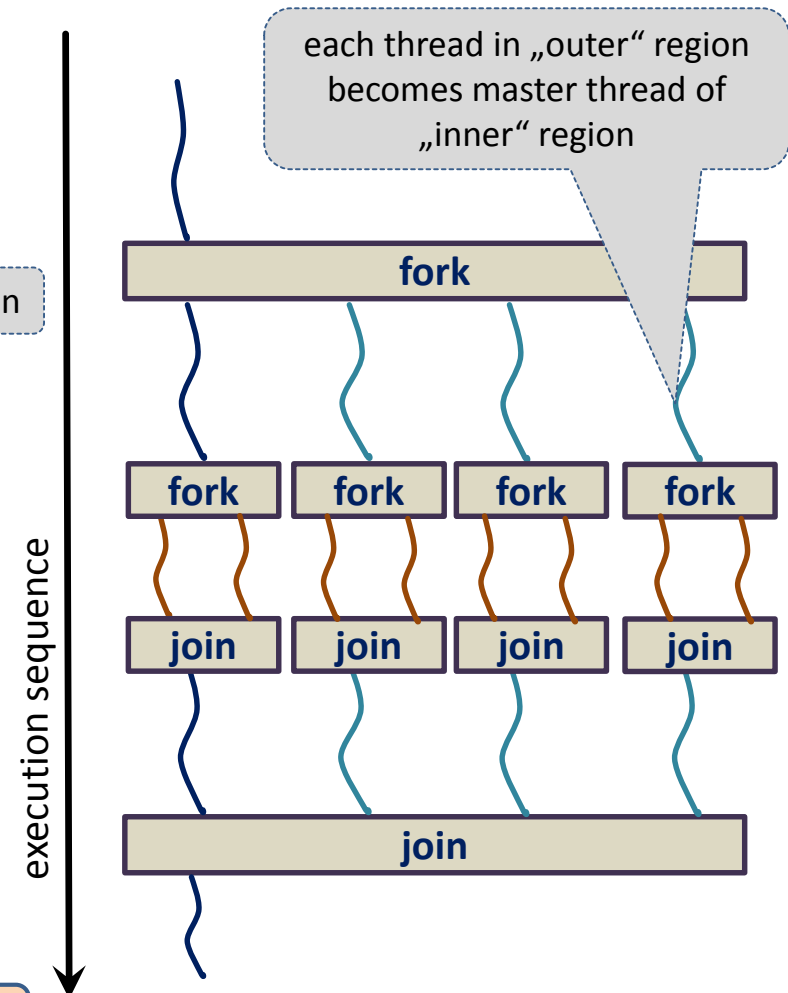
„outer“ region

„inner“ region

C

- nesting of parallel regions

mentioned here for illustrative purposes



■ Suitable environment settings

```
export OMP_NUM_THREADS=4,2  
export OMP_NESTED=true  
export OMP_DYNAMIC=false  
...  
./my_nested_openmp_program.exe
```

one integer for each
nesting level

else, „inner“ regions might/will
execute with 1 thread only.

forbid implementation to interfere with
number of threads assigned

■ Operating system:

- responsible for assigning hardware resources to threads
- in general not trivial – note that (active) thread count can change during execution

■ Possible issues (performance impact):

- threads might move around between cores
- multiple threads might share a core (or other resources)

→ a mechanism for controlling thread affinity / binding is desirable

- **Two aspects:**
 1. What entity should a thread be bound to? → concept of **place**
 2. How should the binding be performed (if at all ...)?

- **Optimal binding strategy** depends on machine and application
- **Putting threads far apart („spread“, „scatter“)** might
 - improve aggregate memory bandwidth
 - improve combined cache size
 - decrease performance of synchronization constructs
- **Putting threads close together** (i.e. on two adjacent cores) might
 - improve performance of synchronization constructs
 - decrease available memory bandwidth and cache size per thread

→ available since **OpenMP 4.0**
before that: implementation-specific mechanisms

■ Places are defined via either

- an abstract name (**threads**, **cores**, or **sockets**), optionally followed by a bracketed positive integer (number of places):

```
export OMP_PLACES="cores(8)"
```

8 places with 1
physical core each

- or an explicit list of places, specified as list of integer intervals (in the following example, all three specs are equivalent)

```
export OMP_PLACES="{0,1,2,3},{4,5,6,7}"
```

2 places with 4 hw
threads each

```
export OMP_PLACES="{0:4},{4:4}"
```

same, using
<offset:length> notation

```
export OMP_PLACES="{0:4}:2:4"
```

same, using
<firstplace:#_of_places:stride_of_offset>
notation

meaning of the index is **implementation defined**, but you can expect the smallest unit of execution (a hardware thread on x86) to be used.

- **Determine whether threads should be pinned**

- environment variable OMP_PROC_BIND
- with values **true** or **false**, or
- a comma-separated list of entries:

master	bind created threads to same place as master thread
close	bind created threads to a place close to the one assigned to the master
spread	use a sparse distribution pattern to bind created threads to places

- **Example:**


```
export OMP_PROC_BIND=spread,close
```

- binding is determined for at most two levels of parallel nesting

- Nested parallelism example from earlier

```
export OMP_NUM_THREADS=4,2
...
export OMP_PLACES="cores(8)"
export OMP_PROC_BIND=spread,close
./my_nested_openmp_program.exe
```

- Threads are named S_i , and S'_i , S''_i , ..., for outer and inner region, respectively:



- Overcommitment causes places to be reused (i.e. multiple threads per place)

- **The function**

```
integer(...) function omp_get_proc_bind()
```

Fortran

```
omp_proc_bind_t omp_get_proc_bind(void)
```

C

returns one of the following **constants**:

omp_proc_bind_false	0
omp_proc_bind_true	1
omp_proc_bind_master	2
omp_proc_bind_close	3
omp_proc_bind_spread	4

- **The value may depend on the nesting level from which the function is called**

- A number of functions exist to handle various inquiries:

Name	Result type	Purpose
<code>omp_get_num_places()</code>	int	number of places available
<code>omp_get_place_num_procs</code> <code>(int place_num)</code>	int	number of processors available in <code>place_num</code> (0 .. number of places - 1)
<code>omp_get_place_proc_ids</code> <code>(int place_num, int *ids)</code>	void	<code>ids</code> contains numerical identifiers of processors in place <code>place_num</code>
<code>omp_get_place_num()</code>	int	place number of place to which calling thread is bound
<code>omp_get_partition_num_places()</code>	int	number of places in place partition of innermost implicit task
<code>omp_get_partition_place_nums</code> <code>(int *place_nums)</code>	void	list of place numbers for innermost implicit task

- A `proc_bind` clause can be specified
- **Example:**


```
#pragma omp parallel num_threads(4) proc_bind(spread)
{ ...
#pragma omp parallel num_threads(2) proc_bind(close)
{ ...
}
}
```

C

executed with
`OMP_PLACES=cores(8)`

Node

outer region
inner region



■ Topology =

- Where in the machine does core #n reside?
- awkward numbering anyway?
- which cores share which cache levels
- which hardware threads (“logical cores”) share a physical core?

■ Use LIKWID tool to identify

- developed by J. Treibig
- see <http://code.google.com/p/likwid> for source code and documentation

■ Available commands

- **likwid-topology**: Print thread and cache topology
- **likwid-pin**: Pin threaded application without touching code
- **likwid-perfctr**: Measure performance counters
- **likwid-mpirun**: mpirun wrapper script for easy LIKWID integration
- **likwid-bench**: Low-level bandwidth benchmark generator tool
- ... some more

■ Output of **likwid-topology -g** (ASCII art section):

Socket 0:

	+-----+			+-----+			+-----+			+-----+			+-----+			+-----+			+-----+		
	0	16		1	17		2	18		3	19		4	20		5	21		6	22	
	+-----+			+-----+			+-----+			+-----+			+-----+			+-----+			+-----+		
	+-----+			+-----+			+-----+			+-----+			+-----+			+-----+			+-----+		
	32kB			32kB			32kB			32kB			32kB			32kB			32kB		
	+-----+			+-----+			+-----+			+-----+			+-----+			+-----+			+-----+		
	+-----+			+-----+			+-----+			+-----+			+-----+			+-----+			+-----+		
	256kB			256kB			256kB			256kB			256kB			256kB			256kB		
	+-----+			+-----+			+-----+			+-----+			+-----+			+-----+			+-----+		
	+-----+																				
	20MB																				
	+-----+																				
+																					

hyperthreaded
cores

L1D

L2

shared
L3

Socket 1:

+-----+-----+-----+-----+-----+-----+-----+-----+															
	+-----+			+-----+			+-----+			+-----+			+-----+		
	8	24		9	25		10	26		11	27		12	28	
	+-----+			+-----+			+-----+			+-----+			+-----+		
	+-----+			+-----+			+-----+			+-----+			+-----+		
	32kB			32kB			32kB			32kB			32kB		
	+-----+			+-----+			+-----+			+-----+			+-----+		
	+-----+			+-----+			+-----+			+-----+			+-----+		
	256kB			256kB			256kB			256kB			256kB		
	+-----+			+-----+			+-----+			+-----+			+-----+		
	+-----+														
	20MB														
	+-----+														
+-----+															

each socket forms
a NUMA domain

■ Output of **likwid-topology -g** (ASCII art section):

Socket 0:

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																								
	0			1			2			3			4			5			6			7		
	64kB			64kB			64kB			64kB			64kB			64kB			64kB			64kB		
	512kB			512kB			512kB			512kB			512kB			512kB			512kB			512kB		
	5MB									5MB														
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																								

single threaded
cores

L1D

L2

shared
L3

Socket 1:

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																								
	8			9			10			11			12			13			14			15		
	64kB			64kB			64kB			64kB			64kB			64kB			64kB			64kB		
	512kB			512kB			512kB			512kB			512kB			512kB			512kB			512kB		
	5MB									5MB														
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																								

each socket forms
two NUMA domains

- **Pins processes/threads to specific cores without touching code**
 - Directly supports pthreads, gcc OpenMP, Intel OpenMP
 - Based on combination of wrapper tool together with overloaded pthread library → binary must be **dynamically linked**!
- **Can also be used as a superior replacement for Linux command `taskset`**
- **Supports logical core numbering within a node and within an existing CPU set**
 - Useful for running inside CPU sets defined by someone else, e.g., the MPI start mechanism or a batch system
- **Usage examples:**
 - Physical numbering (as given by likwid-topology):

```
likwid-pin -c 0,2,4-6 ./myApp parameters
```
 - Logical numbering by topological entities:

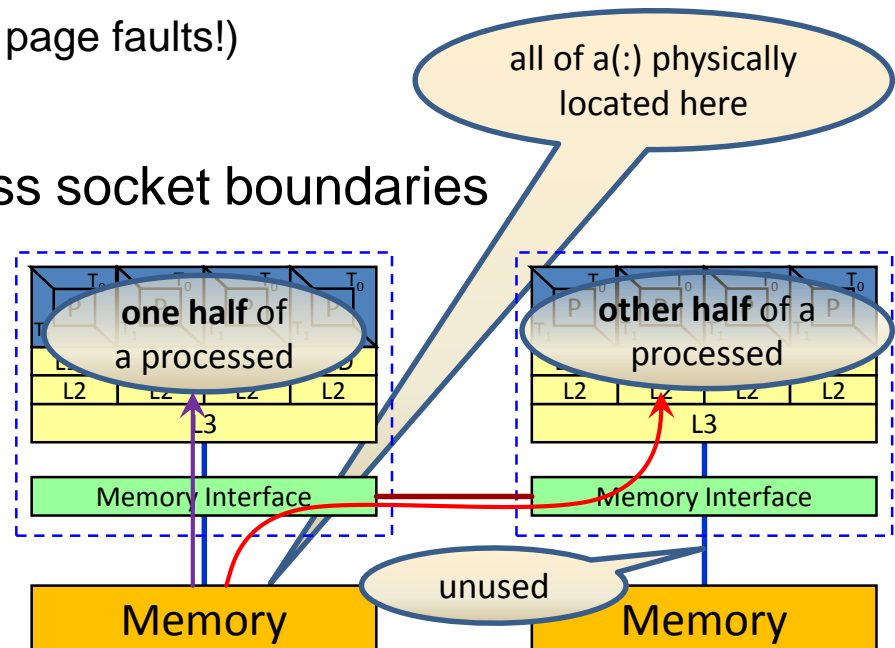
```
likwid-pin -c S0:0-3 ./myApp parameters
```

- **Allocation of memory** (with C malloc() / Fortran ALLOCATE)
 - only provides a virtual memory address
- **Physical memory**
 - is assigned when a memory location is initialized („first touch“)
 - units of pages (note overhead due to page faults!)
- **Consequence for OpenMP**
 - possible memory accesses across socket boundaries

Fortran

```
a(:) = 0.0
!$omp parallel do
DO i=1, size(a)
... = ... a(i) ...
END DO
!$omp end parallel do
```

first touch here

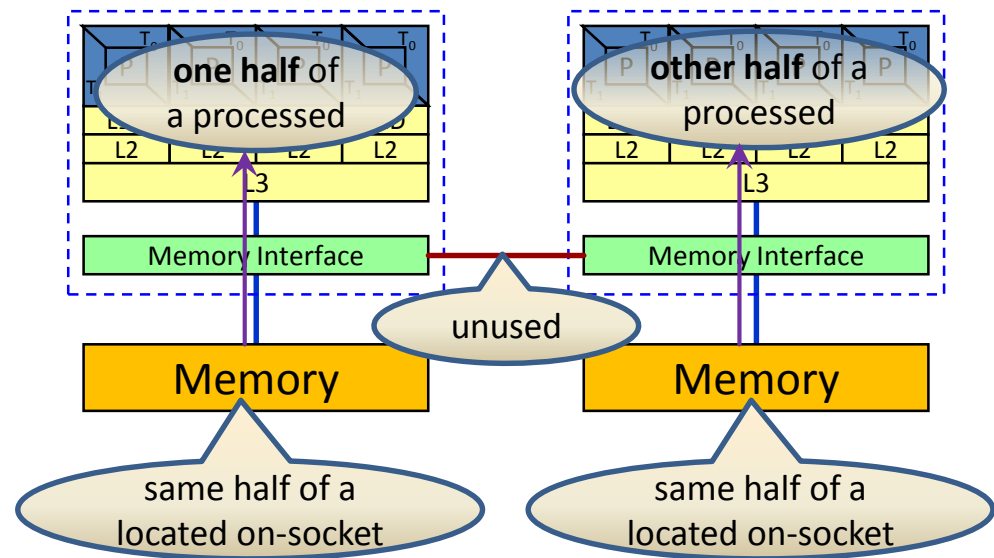


- only **half** the available memory BW might be exploited on a 2-socket system

- **Desirable** and **scalable** memory access pattern:
 - requires initialization with an OpenMP parallelized loop
- **Distributed first touch**
 - ideally, uses same loop schedule as later processing

Fortran


```
!$omp parallel do
DO i=1, size(a)
  a(i) = ...
END DO
!$omp end parallel do
...
!$omp parallel do
DO i=1, size(a)
  ... = ... a(i) ...
END DO
!$omp end parallel do
```



- now, the **full** available memory BW can be exploited on a **multi**-socket system

■ Measured on two AMD Magny Cours sockets

- thread pinning uses „close“ strategy



■ Remember:

- tasking decouples data items and associated functions from the threading model

```
#pragma omp task
  execute_my_function(a, b, c);
```

■ Consequence:


- repeated execution of tasking on data items might use different threads → memory affinity will get lost!

```
#pragma omp task shared(a, b, c)
  establish_my_data(a, b, c);
#pragma omp taskwait
#pragma omp task shared(a, b, c)
  execute_my_function(a, b, c);
```

this function might execute on a **different** thread than this one

■ At initialization

- store which thread performed it – threads are color coded below




```
integer :: work_item(idm, nthr)
```

thread	0	1	2
item #	1	2	3
item #	4	5	6
item #	7	-	-

■ Working on data items

- first work on items that are local to the executing thread
- next work on items that are located elsewhere (nearby first)
 - task stealing due to unpredictable thread assignment
- additional bookkeeping (mutual exclusion) is needed to assure complete and unique execution

work shared vector triad with 16 threads on Sandy Bridge



- Example program: count even and odd array values


```
integer is(2), ict(2,ntdm), ia(n)
...
!$omp parallel private(myid) shared(ict, ia)
  myid = omp_get_thread_num()+1
!$omp do private(index)
  do i=1,n
    index = mod(ia(i),2)+1
    ict(index,myid) = ict(index,myid) + 1
  end do
!$omp end do
!$omp critical
  is = is + ict(1:2,myid)
!$omp end critical
!$omp end parallel
```

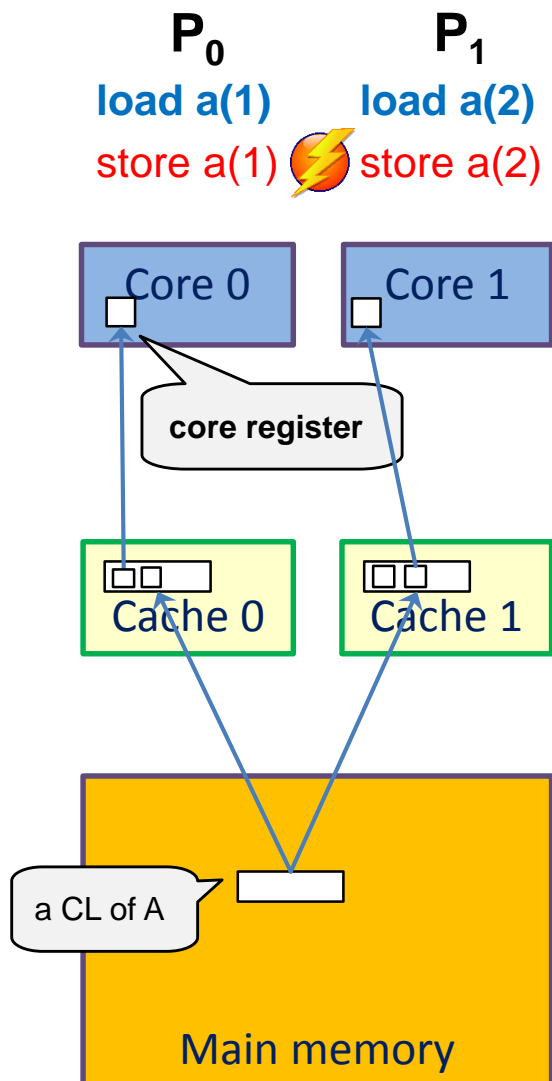
initialization omitted

formally correct,
no race condition

Fortran

- Baseline 1 thread execution time: AMD 0.75 s, Intel SandyBridge 0.37 s





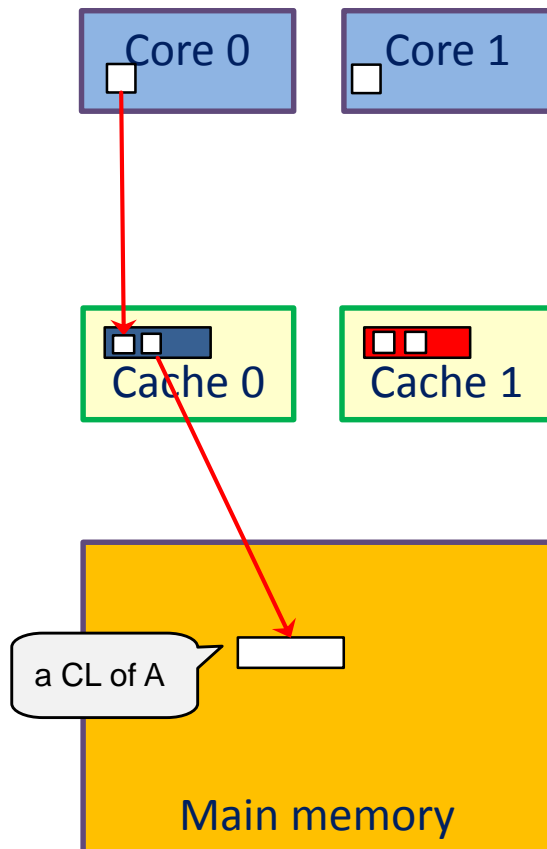
■ Store operation

- write back always done on **complete** cache lines
- "merging of partial cache lines" is **not** possible

■ Cache coherence protocol

- keeps track of cache line status
- assures data consistency by enforcing hardware synchronization between writes

Diagram shows state
after step 3



- **Hardware execution sequence for write on Core 0:**
 1. Request exclusive access to CL (Core 0 issues it first)
 2. **Invalidate** CL in Cache 1
 3. Modify CL in Cache 0 (exclusively owned)
 4. mark CL **shared**
- **Hardware execution sequence on Core 1:**
 5. Request CL from memory for reading (granted after CL is marked shared)
 6. Request exclusive access to CL
 7. **Invalidate** CL in Cache 0
 8. Modify CL in Cache 1 (exclusively owned)
 9. mark CL **shared**

- **Repeated access to data in same cache line:**
 - causes thrashing of cache lines
 - for each access, more than twice the memory latency may be accumulated, resulting in significant performance reduction

- **This effect is called "false sharing"**

- **Privatization – here through use of a reduction variable**

```
integer is(2), ia(n)
```

```
...
```

initialization omitted

```
!$omp parallel shared(ict, ia)
```

```
!$omp do private(index) reduction(+:is)
```

```
  do i=1,n
```

```
    index = mod(ia(i),2)+1
```

```
    is(index) = is(index) + 1
```

```
  end do
```

```
!$omp end do
```


```
!$omp end parallel
```

private variables are assured of using well-separated parts of the physical memory (thread-individual stack or heap)

Fortran

- **Alternative for retaining shared variables: Add padding**
 - tradeoff: may lose spatial locality

- Baseline 1 thread execution time: AMD 0.81 s, Intel SandyBridge 0.36 s



Now: last exercise session



Outlook: Towards quantifying performance


■ Characteristics

- known operation count, load/store count
- some variants of interest:

Kernel	Name	Flops	Loads	Stores
$s = s + a_i * b_i$	Scalar Product	2	2	0
$n^2 = n^2 + a_i * a_i$	Norm	2	1	0
$a_i = b_i * s + c_i$	Linked Triad (Stream)	2	2	1
$a_i = b_i * c_i + d_i$	Vector Triad	2	3	1

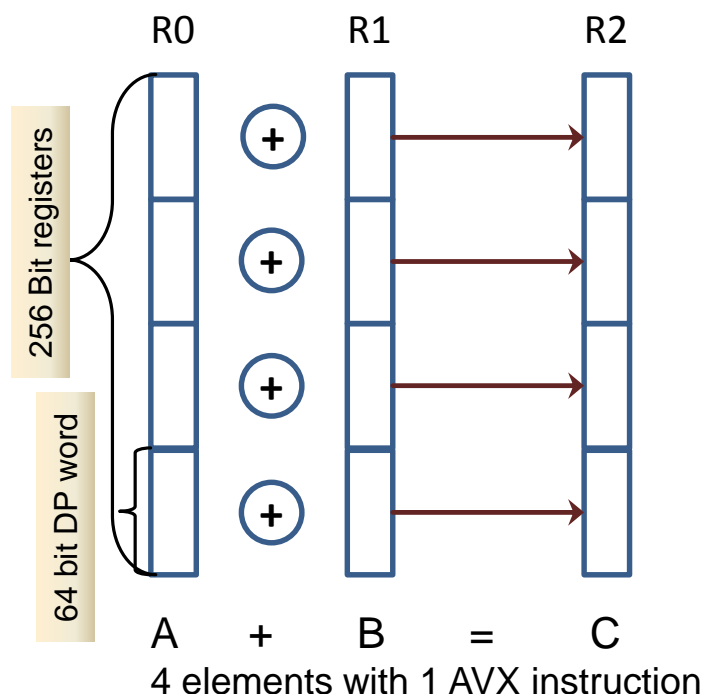
- run repeated iterations for varying vector lengths (working set sizes)

- **Synthetic benchmark:** bandwidths of „raw“ architecture
for a **single core** Sandy Bridge 2.7 GHz / ifort 13.1



■ Sandy Bridge vector unit:

- 256 Bit SIMD (single instruction multiple data)
- Example: addition of 8 Byte words



■ Instruction capability

- 1 vector add and 1 vector mult per cycle → theoretical Peak 8 Flops/cycle

■ LD/ST issue capability

- 4 Words LD/cycle
- 4 Words ST/(2 cycles)


Only L1 might maintain needed bandwidth


■ Vector triad:

- required loads limit performance to 8 Flops / 3 cycles
i.e. **7.2 GFlop/s** at 2.7 GHz

■ Consult processor-specific architecture manual


- **Throughput mode:** run with independent threads **up to number of cores** on a socket





■ Loads and Stores

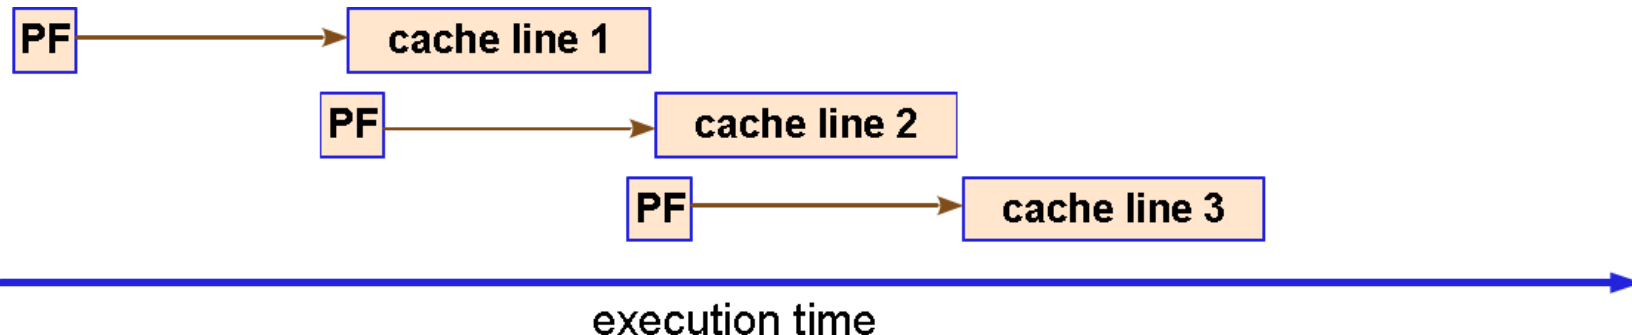
- usually apply to cache lines



- size: 64, 128 or more Bytes

■ Pre-fetch

- avoid latencies when streaming data




- pre-fetches usually done in hardware
- decision according to memory access pattern

■ Pre-Requisite:

- **spatial** locality
- violation of spatial locality:
if only part of a cache line is used
→ effective reduction in bandwidth

$$D(::\text{stride}) = A(::\text{stride}) + B(::\text{stride}) * C(::\text{stride})$$

Example: stride 3




Notes:

- stride known at compile time
- serial compiler optimizations may compensate performance losses in real-life code

← **ca. 40 MFlop/s**
(remains constant for strides > ~25)

■ $\mathbf{r} = \mathbf{M} \cdot \mathbf{x}$ i.e. $r_i = \sum_{j=1}^n M_{ij} x_j$



■ **First parallelization attempt:**

```
!$omp parallel
!$omp do
DO j = 1, n
  DO k = 1, n
    r(j) = r(j) + a(j, k) * x(k)
  END DO
END DO
!$omp end do
... = r(...)
!$omp end parallel
```

index ordering
causes non-contiguous
accesses

■ **Parallel patterns used:**


- data decomposition (load balanced)
- loop parallelism (no dependencies)

■ **Directive placement:**

- coarse grained parallelism to avoid synchronization overhead

- Speed-Up: $S(n_t) = \frac{T(1)}{T(n_t)}$

as a function of number of threads
on 8-core processors




- Scaling **bad** beyond 4 threads

- Absolute performance:

- MFlop/s = $2 \cdot n^2 / \text{time}$

a measure for execution time
if problem size is constant




- used `dgemv` for serial run

- Speed-Up **useless** if baseline performance is bad

- Switch loop order

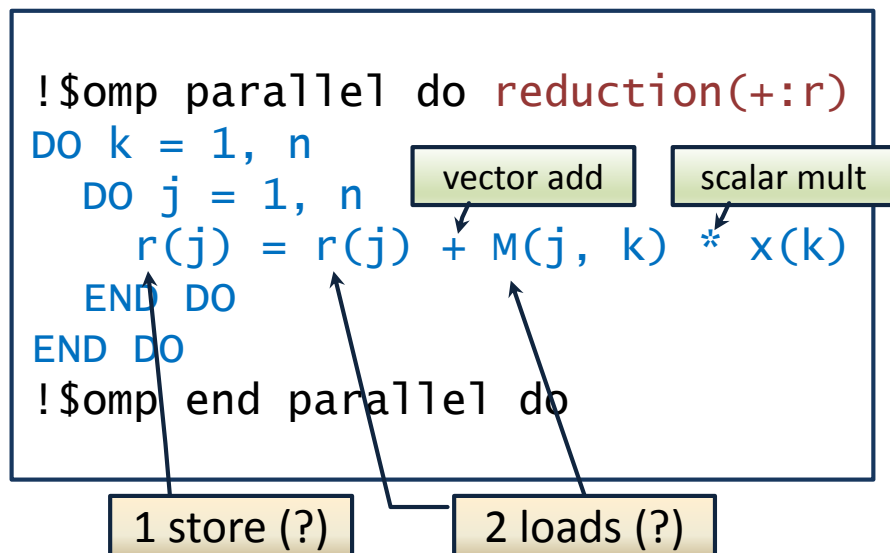
- map **column** blocks to threads:



- color code indicates thread assignment

- Variant 2 of code:

- contiguous** access to M
- array reduction on result vector




- Performance estimate for single thread:

- double that of triad \rightarrow 1.86 GFlop/s

Cannot be the whole truth –
remember serial performance: 3.7 GFlop/s!

- **For variant 2 of the MVM: Performance in MFlop/s**



- **Comments:**

- „no OpenMP“ → variant 2 compiled **without** OpenMP
- Conclusion: compiler stops making certain serial optimizations if OpenMP switch is toggled

Outer loop unrolling

```
!$omp parallel do reduction(+:r)
DO k = 1, n-3, 4
  DO j = 1, n
    r(j) = r(j) + M(j, k) * x(k) &
      + M(j, k+1) * x(k+1) &
      + M(j, k+2) * x(k+2) &
      + M(j, k+3) * x(k+3)
  END DO
END DO
!$omp end parallel do
```

- conditioning omitted
- asymptotically increases intensity to 2 Flops per word (1 load on matrix per original loop iteration)

Unrolling is **limited** by number of available registers and prefetch streams (architecture-dependent!)

Expected performance


- for M from memory (i.e. **outside** any cache)
- contiguous streaming of data
- assuming 40 GB/s bandwidth for a socket

$$\text{Perf} = \frac{2 \text{ Flop}}{8 \text{ Bytes}} * \frac{40 \text{ GB}}{\text{s}} = 10 \text{ GFlop/s}$$


computational
intensity

available
bandwidth
(slowest path)

- estimation method is known as „Roofline Model“



- **In MFlop/s.** Unroll factors: Sandy Bridge 4, Magny Cours 8



- **Comment:**
 - roofline model only predicts „saturated“ performance
 - single-thread performance is limited by non-overlapping memory/core operations (see ref. (2))


- ... if variant 2 gives us the full performance anyway?
 - even if this only is attained with 8 threads

- **Possible reasons:**
 - „switch off“ cores 6-8 to save energy (relevant for you if this is budgeted – may happen not too far in the future!)
 - use cores 6-8 for other tasks that are cache bound
 - use cores 6-8 for MPI communication (I/O via PCI) if you do hybrid programming (i.e., combine MPI with OpenMP)



References

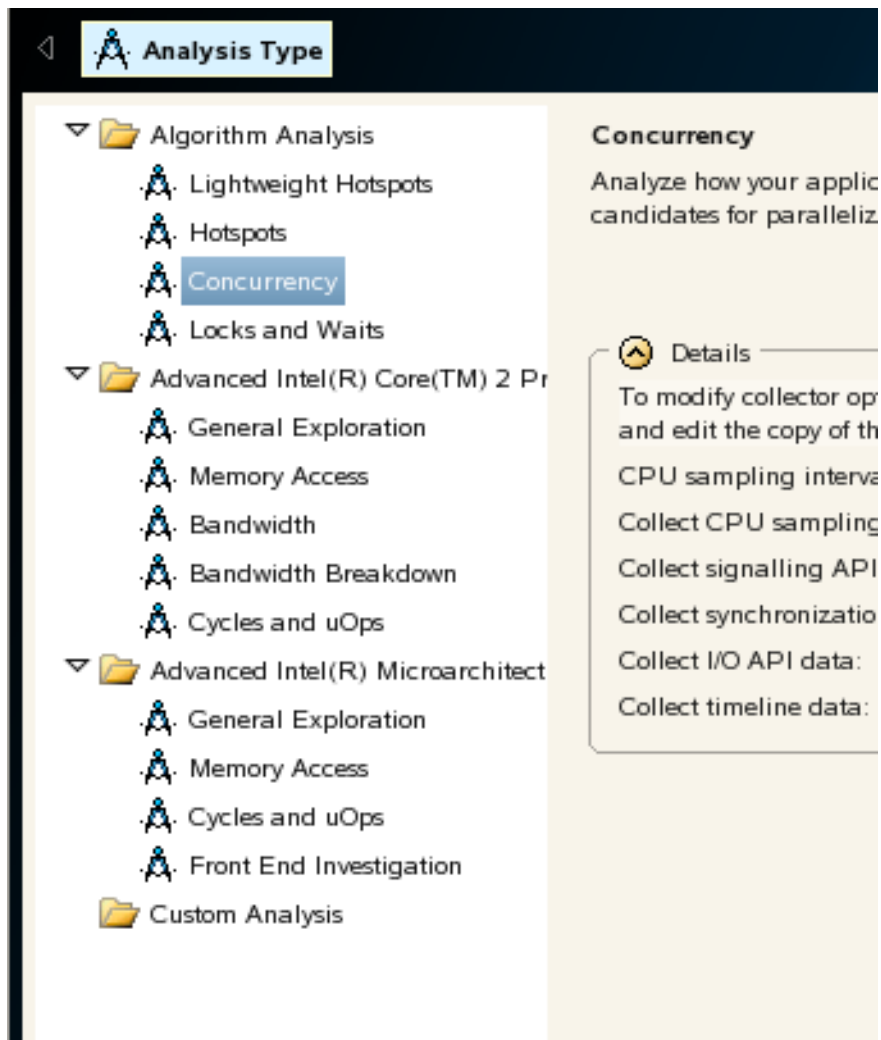
- (1) **OpenMP 4.5 standard and examples (currently 4.0.2)** at
<http://openmp.org/wp/openmp-specifications/>
- (2) **Parallel programming in OpenMP**
Rohit Chandra et al; Morgan Kaufmann 2000
- (3) **Using OpenMP - portable shared memory parallel programming**
B. Chapman, G. Jost, R. van der Pas; MIT Press 2008
- (4) **J. Treibig, G. Hager, G. Wellein: LIKWID**
A lightweight performance-oriented tool suite for x86 multicore environments.
PSTI2010, Sep 13-16, 2010, San Diego, CA DOI: 0.1109/ICPPW.2010.38;
Preprint: <http://arxiv.org/abs/1004.4431>
- (5) **G. Hager, J. Treibig, J. Habich, and G. Wellein:**
Exploring performance and power properties of modern multicore chips via
simple machine models. Preprint: arXiv:1208.2908
- (6) **G. Hager, G. Wellein:** Introduction to High Performance Computing for
Scientists and Engineers. Chapman & Hall / CRC (2011)



Appendix: Setting up Vtune Amplifier

- **Tuning of serial and threaded programs**
 - performance counter access requires group rights
- **Start up GUI**
 - prerequisites: set up environment and possibly stack limit
 - then, invoke the GUI with `amplxe-gui &`
 - command line `amplxe-cl` is also available, but will not be discussed
- **Project generation analogous to Intel Inspector**

```
#pragma omp parallel private(seed,i,k,me)
{
    me = omp_get_thread_num();
    seed = 123 + 159*me;
    for (k=0; k<100000; ++k) {
#pragma omp for
        for (i=0; i<10000; ++i) {
            ir[i] = rand_r(&seed) & 0xf;
        }
#pragma omp master
        for (i=0;i<10000; ++i) {
            hist[ir[i]]++;
        }
#pragma omp barrier
        // prevents ir from being modified
        // before hist update is done
    }
}
```

■ **Various types are provided**

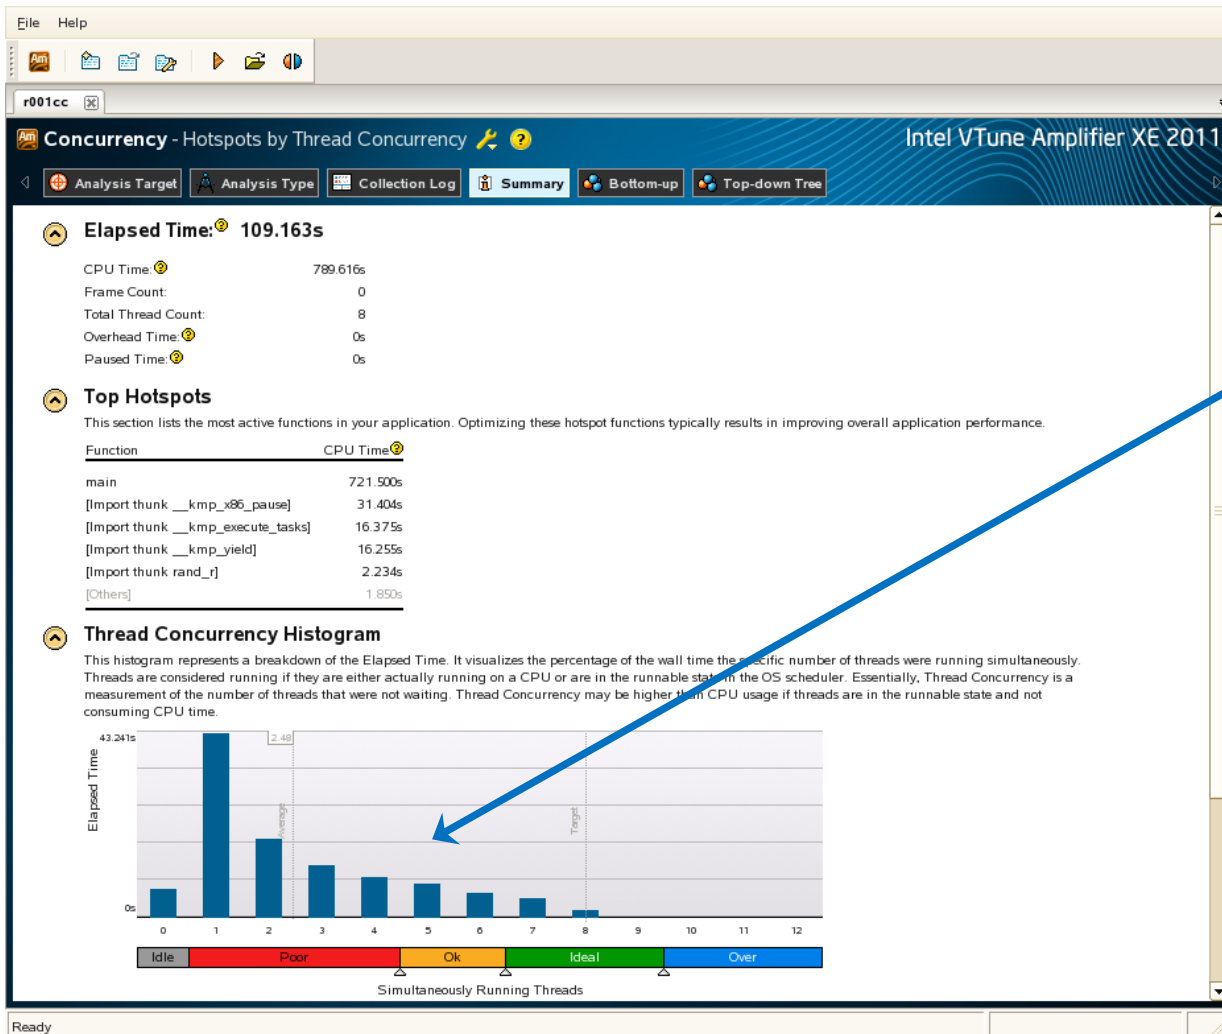
- select „Concurrency“
- in the project properties, set OMP_NUM_THREADS to number of physical cores

Note:

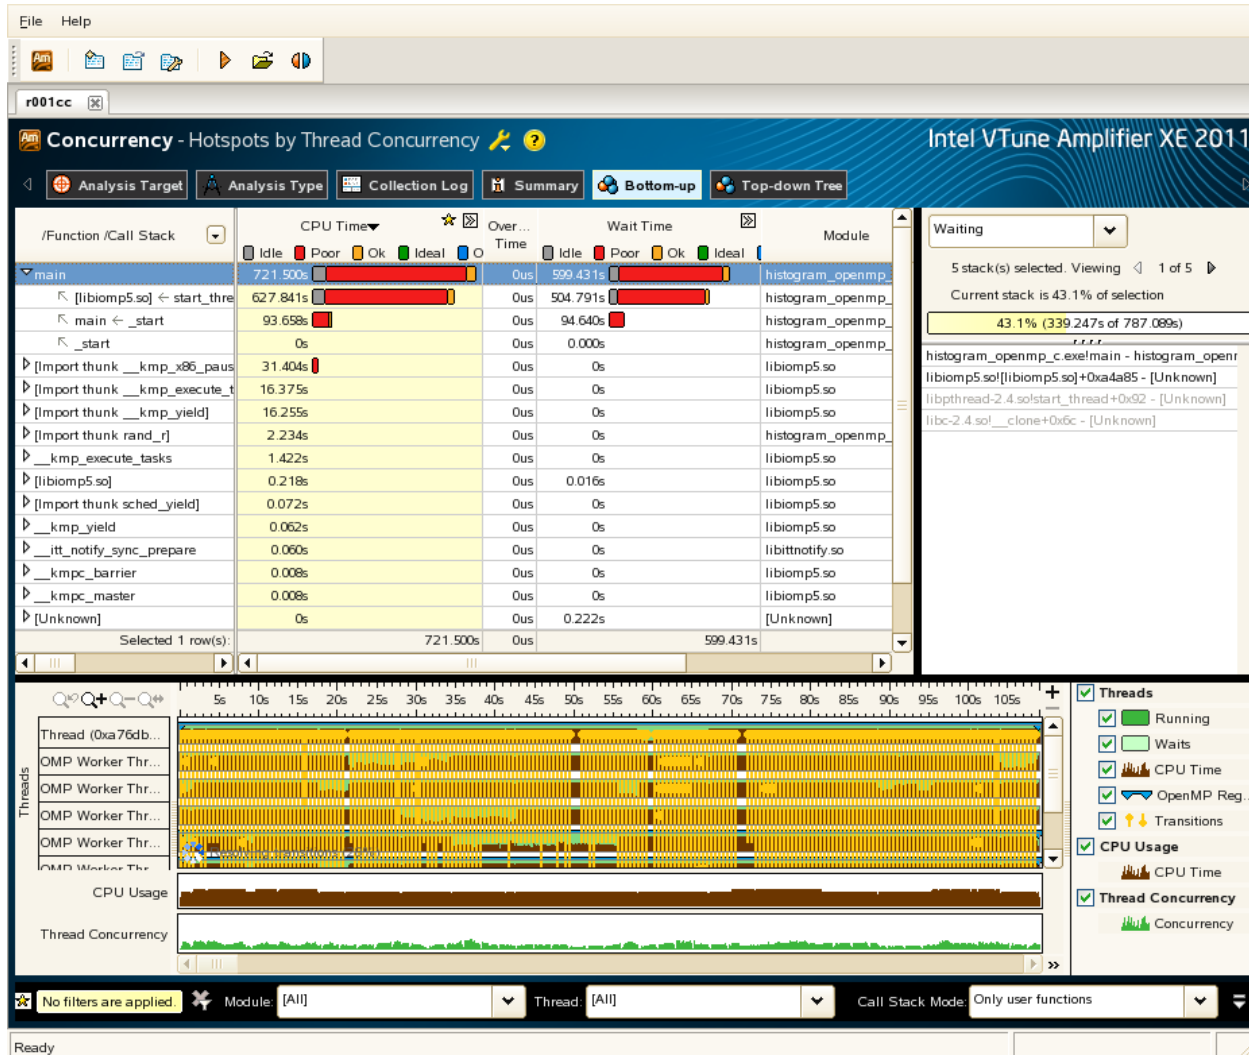
performance quality evaluation assumes complete system is used

■ **Note:**

- analysis may take quite a long time to run, even for programs of small size



- **Result:**
 - thread concurrency very low although CPU usage is high



Observation:

- much time spent in OpenMP run time library
- lots of transitions indicated → have false sharing

- Click on routine with significant resource usage

